

FLUME INSIGHT

WELCOME
“We Will Open the World of
Creativity for you !”




Presented by
Siva Kumar Bhuchipalli



Contents

- Overview of Flume
- Flume Architecture



Flume

Flume is a **highly reliable, distributed** and configurable streaming **data collection tool**.

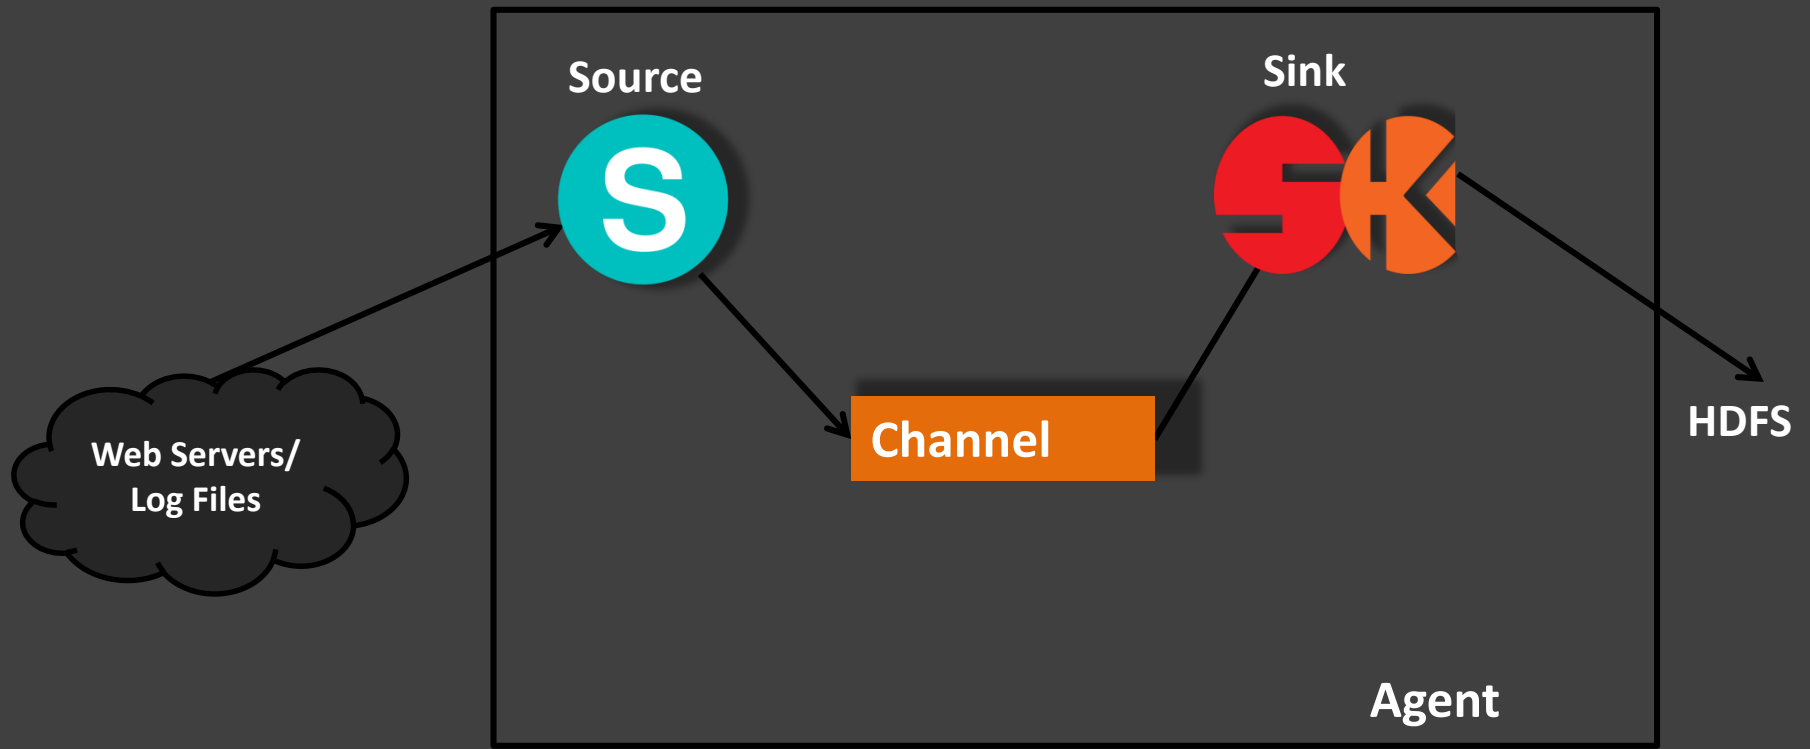
Flume can transport log files across a large number of hosts into HDFS.



Flume Features

- ❖ Flume **collects data** efficiently, **aggregates** and moves large amounts of log data from many different sources to a centralized data store.
- ❖ **Simple and flexible architecture**, and provides a streaming of data flows and leverages data movement from multiple machines in within an enterprise into Hadoop.
- ❖ Flume is **not restricted to log data** aggregation and **it can transport** massive quantities of **event data** including but not limited to network traffic data, social-media-generated data, email messages and pretty much any data source possible.
- ❖ Built-in support for several Sources and destination platforms to integrate with.
- ❖ Support for **multi-hop flows** where events travel through multiple agents before reaching the final destination. fan-in and fan-outflows, contextual routing and backup routes for failed hops are also allowed.

Flume Architecture



Source

- ❖ Flume source is configured within an agent and it **listens for events** from an external source (eg: web server) it reads data, translates events, and handles failure situations.
- ❖ But source doesn't know how to store the event. So, after receiving enough data to produce a Flume event, it **sends events** to the **channel** to which the source is connected.
- ❖ The external source sends events to Flume in a format that is recognized by the target Flume source.
 - ☐ Spooling Directory Source
 - ☐ Exec Source
 - ☐ Netcat Source
 - ☐ HTTP Source
 - ☐ Twitter Source
 - ☐ Avro Source



Channel

- ❖ Channels are **communication bridges** between sources and sinks within an agent. Once a Flume source receives an event, it stores it into one or more channels. The channel is a **passive store that keeps the event until it's consumed by a Flume sink**.
- ❖ **Memory channel** stores the events from in an in-memory queue and from there events will be accessed by sink. Because of software or hardware failures, if the agent process dies in the middle, then all the events currently in the memory channel are lost forever.
- ❖ The **File channel** is another example – it is backed by the local file system. Unlike memory channel, file channel writes the contents to a file on the file system that is deleted only after successful delivery to the sink.
 - ❑ Memory Channel
 - ❑ File Channel
 - ❑ Spillable Memory Channel

Note:

The **memory channel** is the **fastest** but has the **risk of data loss**. The **file channels** are typically much **slower** but effectively provide **guaranteed delivery** to the sink.

Channel

Sink

- ❖ Sink **removes the event from the channel** and **puts it into** an external repository like **HDFS** or forwards it to the Flume source of the next Flume agent in the flow. The source and sink within the given agent run asynchronously with the events staged in the channel.

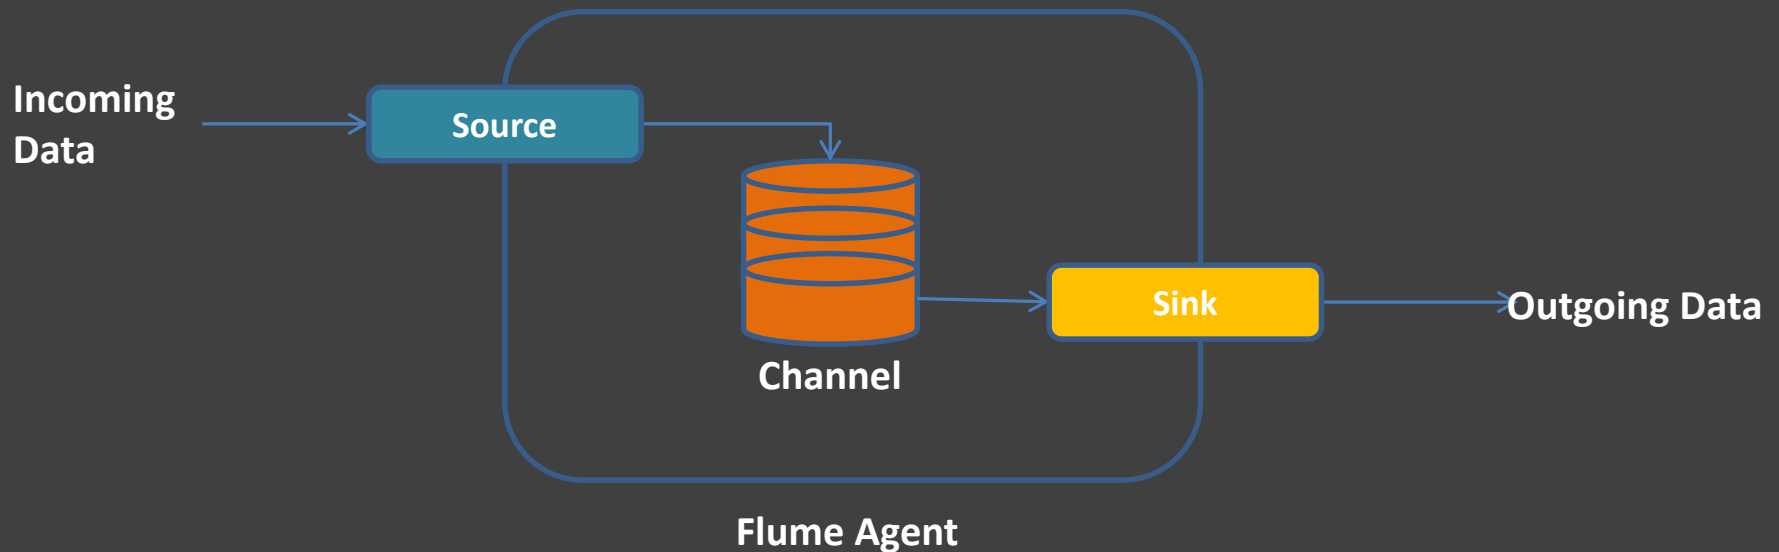
- ☐ HDFS Sink
- ☐ Logger Sink
- ☐ File Roll Sink
- ☐ HBaseSink
- ☐ MorphlineSolrSink
- ☐ ElasticSearchSink
- ☐ Avro Sink



Agent

A container for hosting Sources, Channels, Sinks and other components that enable the transportation of events from one place to another.

- ❖ Fundamental part of a Flume flow
- ❖ Provides Configuration, Life-Cycle Management, and Monitoring Support for hosted components



Anatomy of flume Agent

Aggregation Use Cases.....




Starting Out Simple

- ❖ You would like to move your web server logs to HDFS
- ❖ Let's assume there are only 3 web servers at the time of launch
- ❖ Ad-hoc solution will likely suffice!

Challenges

- ❖ How do you manage your output paths on HDFS?
- ❖ How do you maintain your client code in face of changing environment as well as requirements?




Adding Simple Flume Agent

Advantages

- ❖ Quick offload of logs from Web Server machines
- ❖ Insulation from HDFS downtime
- ❖ Better Network utilization

Challenges

- ❖ What if the Flume node goes down?
- ❖ Can one Flume node accommodate all load from Web Servers?




Adding Two Flume Agent

Advantages

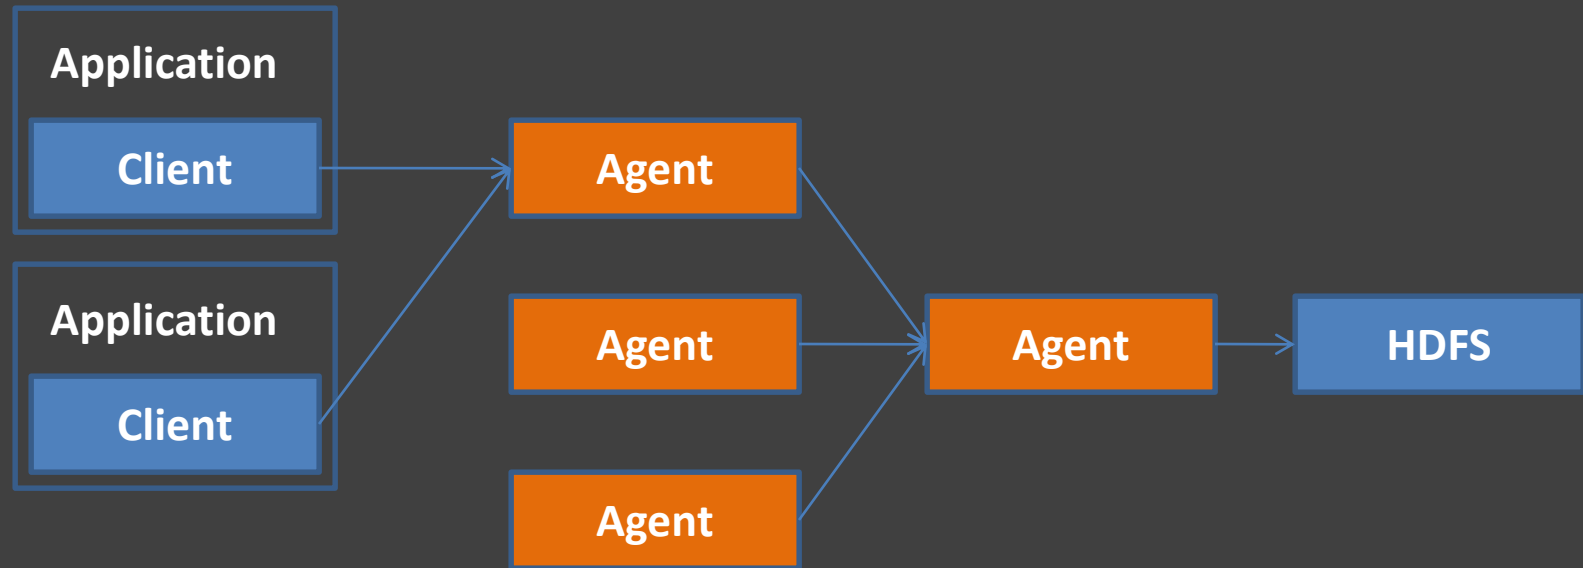
- ❖ Redundancy and Availability
- ❖ Better handling of downstream failures
- ❖ Automatic load balancing and failover

Challenges

- ❖ What happens when new Web Servers are added?
- ❖ Can two Flume Agents keep up with all the load from more Web Servers?




Typical Converging Flow



[Client] → Agent[→ Agent]* → Destination 24

Handling Server Farm

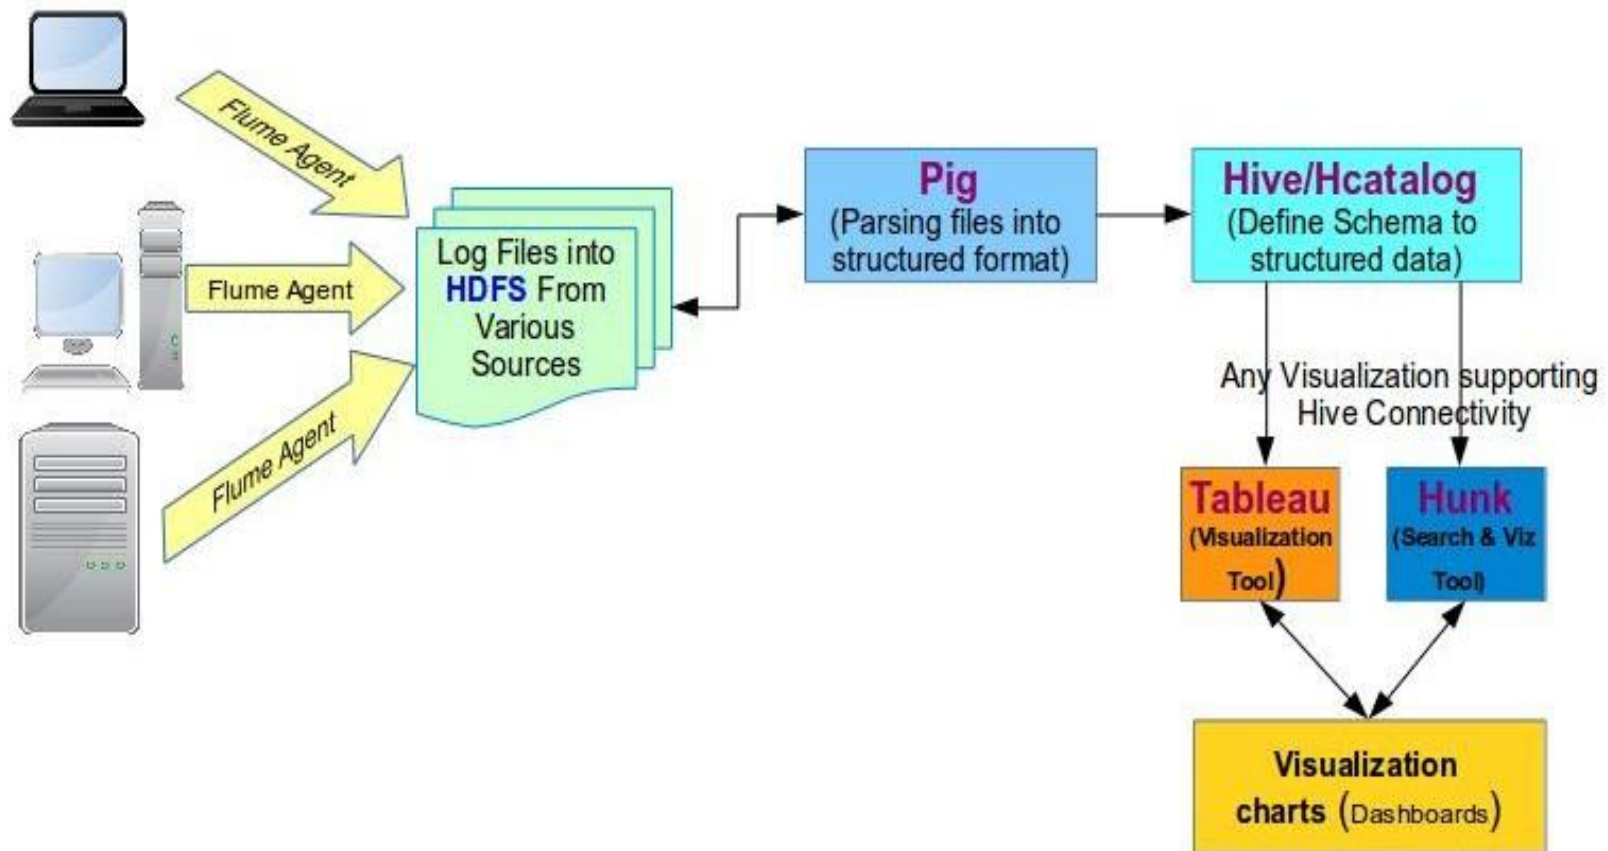


A Converging Flow

- ❖ Traffic is aggregated by **Tier-2** and **Tier-3** before being put into destination system
- ❖ Closer a tier is to the destination, larger the **batch size** it delivers Downstream
- ❖ Optimized handling of destination systems

Log Analysis in Hadoop

Log File Analysis in Hadoop Eco System - Architecture



Setup Of Flume Agent

❖ Flume Agent Setup:

- ☐ Step 1 – Create flume.conf file
- ☐ Step 2 – Starting Flume Agent
- ☐ Step 3 – Feeding events to Flume Agent
- ☐ Step 4 – Verify the events at Agent Console
- ☐ Step 5 – Stop Flume Agent

Sample Flume.conf - NetcatSource, Log Sink and Memory Channel

```
3 # Name the components on this agent
4 Agent1.sources = netcat-source
5 Agent1.channels = memory-channel
6 Agent1.sinks = logger-sink
7
8 # Describe/configure Source
9 Agent1.sources.netcat-source.type = netcat
10 Agent1.sources.netcat-source.bind = localhost
11 Agent1.sources.netcat-source.port = 44444
12
13 # Describe the sink
14 Agent1.sinks.logger-sink.type = logger
15
16 # Use a channel which buffers events in memory
17 Agent1.channels.memory-channel.type = memory
18 Agent1.channels.memory-channel.capacity = 1000
19 Agent1.channels.memory-channel.transactionCapacity = 100
20
21 # Bind the source and sink to the channel
22 Agent1.sources.netcat-source.channels = memory-channel
23 Agent1.sinks.logger-sink.channel = memory-channel
```

Start Flume Agent

```
$ flume-ng agent --conf flume/conf/path --conf-file flume/conf/path/flume.conf --  
name Agent1 -Dflume.root.logger=INFO,console
```

Feeding Events

```
$ curl telnet://localhost:44444
```

or

```
$ telnet localhost 44444
```

Agent – SpoolingDir Source, File Channel, HDFS Sink

```
agent1.sinks = hdfs-sink1_1  
agent1.sources = source1_1  
agent1.channels = fileChannel1_1
```

```
agent1.sources.source1_1.type = spoolDir  
agent1.sources.source1_1.spoolDir = /home/cloudera/spool  
agent1.sources.source1_1.fileHeader = false  
agent1.sources.source1_1.fileSuffix = .COMPLETED
```

```
agent1.sinks.hdfs-sink1_1.type = hdfs  
agent1.sinks.hdfs-sink1_1.hdfs.path = hdfs://quickstart.cloudera:8020/flume_sink/%y-%m-%d/%H%M/  
agent1.sinks.hdfs-sink1_1.hdfs.batchSize = 1000  
agent1.sinks.hdfs-sink1_1.hdfs.rollSize = 268435456  
agent1.sinks.hdfs-sink1_1.hdfs.rollInterval = 0  
agent1.sinks.hdfs-sink1_1.hdfs.rollCount = 50000000  
agent1.sinks.hdfs-sink1_1.hdfs.writeFormat=Text  
agent1.sinks.hdfs-sink1_1.hdfs.useLocalTimeStamp = true  
agent1.sinks.hdfs-sink1_1.hdfs.fileType = DataStream
```

```
agent1.sources.source1_1.channels = fileChannel1_1  
agent1.sinks.hdfs-sink1_1.channel = fileChannel1_1
```

```
agent1.channels.fileChannel1_1.type = file  
agent1.channels.fileChannel1_1.checkpointDir = /home/cloudera/flume/checkpoint/  
agent1.channels.fileChannel1_1.dataDirs = /home/cloudera/flume/data/
```

HDFS Sink Properties

| Property Name | Default | Description |
|-----------------------|--------------|---|
| type | – | The component type name, needs to be hdfs |
| hdfs.path | – | HDFS directory path (eg hdfs://namenode/flume/webdata/) |
| hdfs.filePrefix | FlumeData | Name prefixed to files created by Flume in hdfs directory |
| hdfs.fileSuffix | – | Suffix to append to file (eg .avro – NOTE: period is not automatically added) |
| hdfs.inUseSuffix | .tmp | Suffix that is used for temporal files that flume actively writes into |
| hdfs.rollInterval | 30 | Number of seconds to wait before rolling current file (0 = never roll based on time interval) |
| hdfs.rollSize | 1024 | File size to trigger roll, in bytes (0: never roll based on file size) |
| hdfs.rollCount | 10 | Number of events written to file before it rolled (0 = never roll based on number of events) |
| hdfs.batchSize | 100 | number of events written to file before it is flushed to HDFS |
| hdfs.codeC | – | Compression codec. one of following : gzip, bzip2, lzo, lzop, snappy |
| hdfs.fileType | SequenceFile | File format: currently SequenceFile , DataStream or CompressedStream (1)DataStream will not compress output file and please don't set codeC (2)CompressedStream requires set hdfs.codeC with an available codeC |
| hdfs.minBlockReplicas | – | Specify minimum number of replicas per HDFS block. If not specified, it comes from the default Hadoop config in the classpath. |
| serializer | TEXT | Other possible options include avro_event or the fully-qualified class name of an implementation of the EventSerializer.Builder interface. |

Twitter Agent

TwitterAgent.sources = Twitter

TwitterAgent.channels = FileChannel

TwitterAgent.sinks = HDFS

#TwitterAgent.sources.Twitter.type = com.cloudera.flume.source.TwitterSource

TwitterAgent.sources.Twitter.type = org.apache.flume.source.twitter.TwitterSource

TwitterAgent.sources.Twitter.channels = FileChannel

TwitterAgent.sources.Twitter.consumerKey = RUez5WLB8UxhgfhdsdjVQTs

TwitterAgent.sources.Twitter.consumerSecret = gZN2xtwoRWeruB9ghsgssfk201Pt3tgEBRPWH6feML5oq6ho

TwitterAgent.sources.Twitter.accessToken = 56093598-ALvkeeLbhzQwBoDgO1298sfhj98nfa51BpgANUxTt

TwitterAgent.sources.Twitter.accessTokenSecret = OtfiCBSXjiApSSmZk70h31jkajjadkaXoAugnC9pSm

TwitterAgent.sources.Twitter.maxBatchSize = 50000

TwitterAgent.sources.Twitter.maxBatchDurationMillis = 100000

TwitterAgent.sinks.HDFS.channel = FileChannel

TwitterAgent.sinks.HDFS.type = hdfs

TwitterAgent.sinks.HDFS.hdfs.path = hdfs://localhost:9000/user/flume/tweets/

TwitterAgent.sinks.HDFS.hdfs.fileType = DataStream

TwitterAgent.sinks.HDFS.hdfs.writeFormat = Text

TwitterAgent.sinks.HDFS.hdfs.batchSize = 200000

TwitterAgent.sinks.HDFS.hdfs.rollSize = 0

TwitterAgent.sinks.HDFS.hdfs.rollCount = 2000000

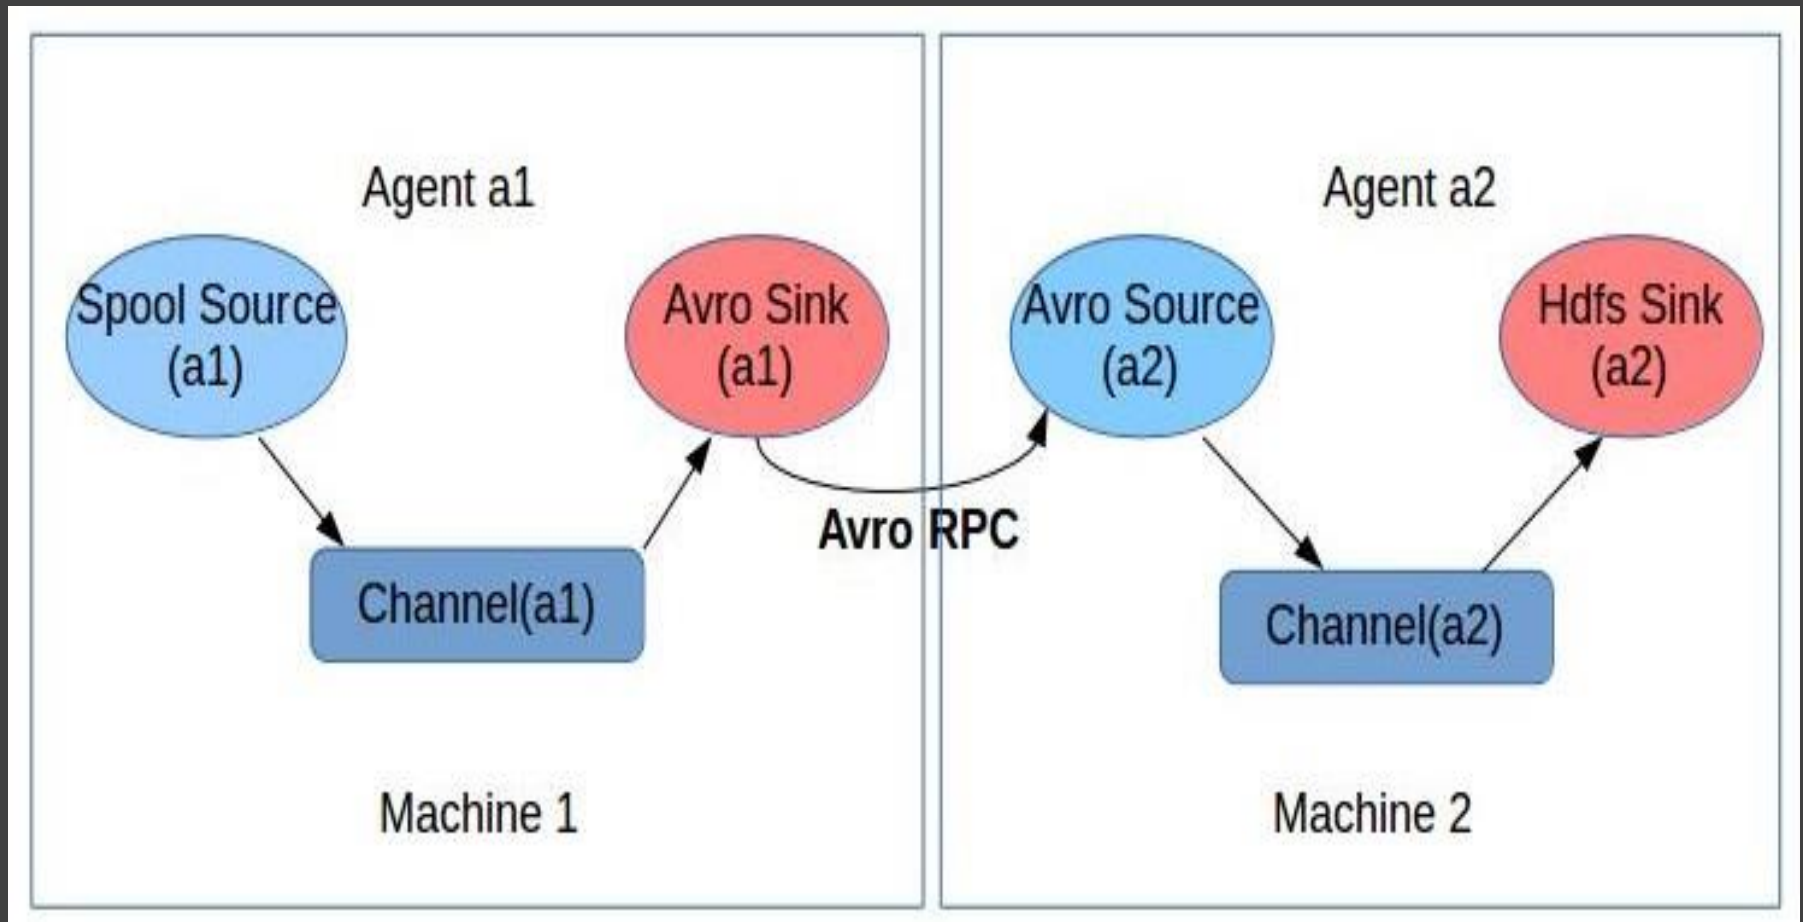
TwitterAgent.channels.FileChannel.type = file

TwitterAgent.channels.FileChannel.checkpointDir = /var/log/flume/checkpoint/

TwitterAgent.channels.FileChannel.dataDirs = /var/log/flume/data/

<http://hadooptutorial.info/twitter-data-analysis-using-hadoop-flume/>

Multi Agent Flow



Multi Agent Flow

```
### Agent1 - Spooling Directory Source and File Channel, Avro Sink ###
```

```
# Name the components on this agent
```

```
Agent1.sources = spooldir-source
```

```
Agent1.channels = file-channel
```

```
Agent1.sinks = avro-sink
```

```
# Describe/configure Source
```

```
Agent1.sources.spooldir-source.type = spooldir
```

```
Agent1.sources.spooldir-source.spoolDir = /home/user/testflume/spooldir
```

```
# Describe the sink
```

```
Agent1.sinks.avro-sink.type = avro
```

```
Agent1.sinks.avro-sink.hostname = 251.16.12.112 #IP Address masked here
```

```
Agent1.sinks.avro-sink.port = 11111
```

```
#Use a channel which buffers events in file
```

```
Agent1.channels.file-channel.type = file
```

```
Agent1.channels.file-channel.checkpointDir = /home/user/testflume/checkpoint/
```

```
Agent1.channels.file-channel.dataDirs = /home/user/testflume/data/
```

```
# Bind the source and sink to the channel
```

```
Agent1.sources.spooldir-source.channels = file-channel
```

```
Agent1.sinks.avro-sink.channel = file-channel
```

<http://hadooptutorial.info/multi-agent-setup-in-flume/>

Multi Agent Flow

```
### Agent2 - Avro Source and File Channel, Avro Sink ###  
# Name the components on this agent  
Agent2.sources = avro-source  
Agent2.channels = file-channel  
Agent2.sinks = hdfs-sink  
  
# Describe/configure Source  
Agent2.sources.avro-source.type = avro  
Agent2.sources.avro-source.hostname = 251.16.12.112  
Agent2.sources.avro-source.port = 11111  
  
# Describe the sink  
Agent2.sinks.hdfs-sink.type = hdfs  
Agent2.sinks.hdfs-sink.hdfs.path = hdfs://251.16.12.112:9000/flume/  
Agent2.sinks.hdfs-sink.hdfs.rollInterval = 0  
Agent2.sinks.hdfs-sink.hdfs.rollSize = 0  
Agent2.sinks.hdfs-sink.hdfs.rollCount = 10000  
Agent2.sinks.hdfs-sink.hdfs.fileType = DataStream  
  
#Use a channel which buffers events in file  
Agent2.channels.file-channel.type = file  
Agent2.channels.file-channel.checkpointDir = /home/user/testflume/checkpoint/  
Agent2.channels.file-channel.dataDirs = /home/user/testflume/data/  
  
# Bind the source and sink to the channel  
Agent2.sources.avro-source.channels = file-channel  
Agent2.sinks.hdfs-sink.channel = file-channel
```

Some Highlights Of Flume

- ❖ Flume is suitable for large volume data collection, especially when data is being produced in multiple locations
- ❖ Once planned and sized appropriately, Flume will practically run itself without any operational intervention
- ❖ Flume provides weak ordering guarantee, i.e., in the absence of failures the data will arrive in the order it was received in the Flume pipeline
- ❖ Flume is a continuous service that will be monitoring arrival of input and ingests that into destination without our directions to start the process every time data arrives.
- ❖ Flume has rich out-of-the box features such as contextual routing, and support for popular data sources and destination systems

QUESTIONS



THANK
YOU

