



# HBASE



Presented By  
Siva Kumar Bhuchipalli

Your **FUTURE** is created by what

**YOU** Do

~~Not~~

**TODAY**

~~**TOMORROW**~~

- **Column-Oriented data store**, known as “**Hadoop Database**”
- Supports random real-time CRUD operations (unlike HDFS)
- Distributed – designed to serve large tables
  - ✓ Billions of rows and millions of columns
- Runs on a cluster of commodity hardware
  - ✓ Server hardware, not laptop/desktops
- Open-source, written in Java
- Type of “**NoSQL**” DB
  - ✓ Does not provide a SQL based access
  - ✓ Does not adhere to Relational Model for storage
- Horizontally scalable
  - ✓ **Automatic sharding\***
- Strongly consistent reads and writes
- Automatic fail-over
- Simple Java API
- Integration with Map/Reduce framework
- Thrift, Avro and REST-ful Web-services



# When To Use HBASE?

- Not suitable for every problem
  - ✓ Compared to **RDBMs** has VERY simple and limited API
- Good for large amounts of data
  - ✓ 100s of millions or billions of rows
  - ✓ If data is too small all the records will end up on a single node leaving the rest of the cluster idle

- HBase is memory and CPU intensive



- Two well-known use cases
  - ✓ Lots and lots of data (already mentioned)
  - ✓ Large amount of clients/requests (usually cause a lot of data)
- Great for single random selects and range scans by key
- Great for variable schema
  - ✓ Rows may drastically differ
  - ✓ If your schema has many columns and most of them are null



# HBASE Data Model?

- Namespaces are Databases

Data is stored in **Tables**

- Tables contain rows
  - ✓ Rows are referenced by a **unique key**
- Key is an array of bytes – good news
- Anything can be a key: string, long and your own serialized data structures
- Rows made of columns which are grouped in column families
- Data is stored in cells
  - ✓ Identified by row x **column-family** x column
  - ✓ Cell's content is also an array of bytes




# HBASE Families?

- Columns are grouped into families
  - ✓ Labeled as “**family:column**”
- Example “**user:first\_name**”
  - ✓ A way to organize your data
- Compression
- In-memory option
- Stored together - in a file called HFile/StoreFile
- Family definitions are static initially, it is no longer static
  - ✓ Created with table, should be rarely added and changed
  - ✓ Limited to small number of families
- unlike columns that you can have millions of
- Family name must be composed of printable characters
  - ✓ Not bytes, unlike keys and values
- Think of **family:column** as a tag for a cell value and NOT as a spreadsheet
- Columns on the other hand are NOT static
  - ✓ Create new columns at run-time
  - ✓ Can scale to millions for a family

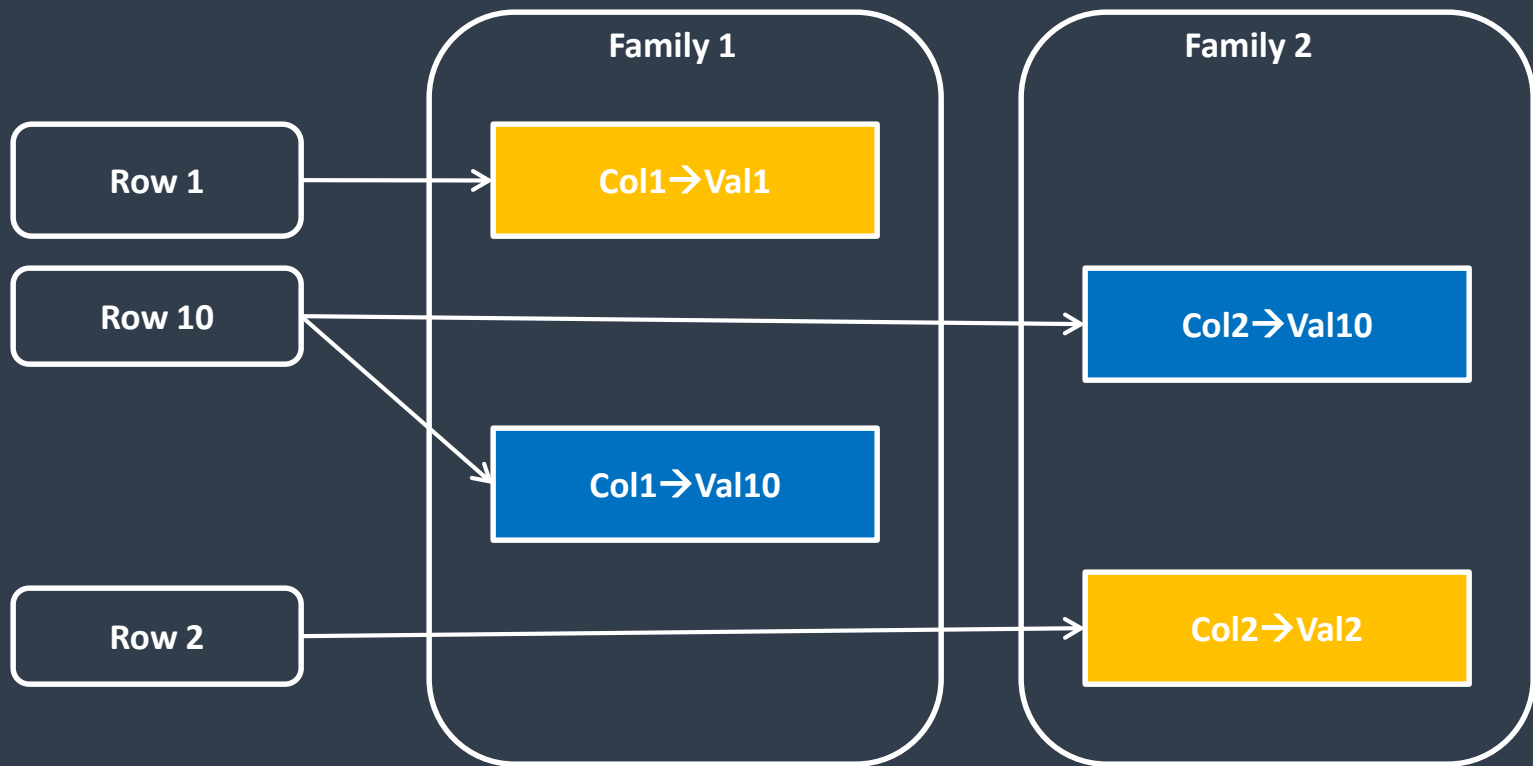


# HBASE Row Keys

- Rows are sorted lexicographically by key
  - ✓ Compared on a binary level from left to right
  - ✓ For example keys 1,2,3,10,15 will get sorted as 1, 10, 15, 2, 3
- Somewhat similar to Relational DB primary index
  - ✓ Always **unique**
  - ✓ Some but minimal secondary indexes support



# Row Composed Of Cells Stored In Families : Columns





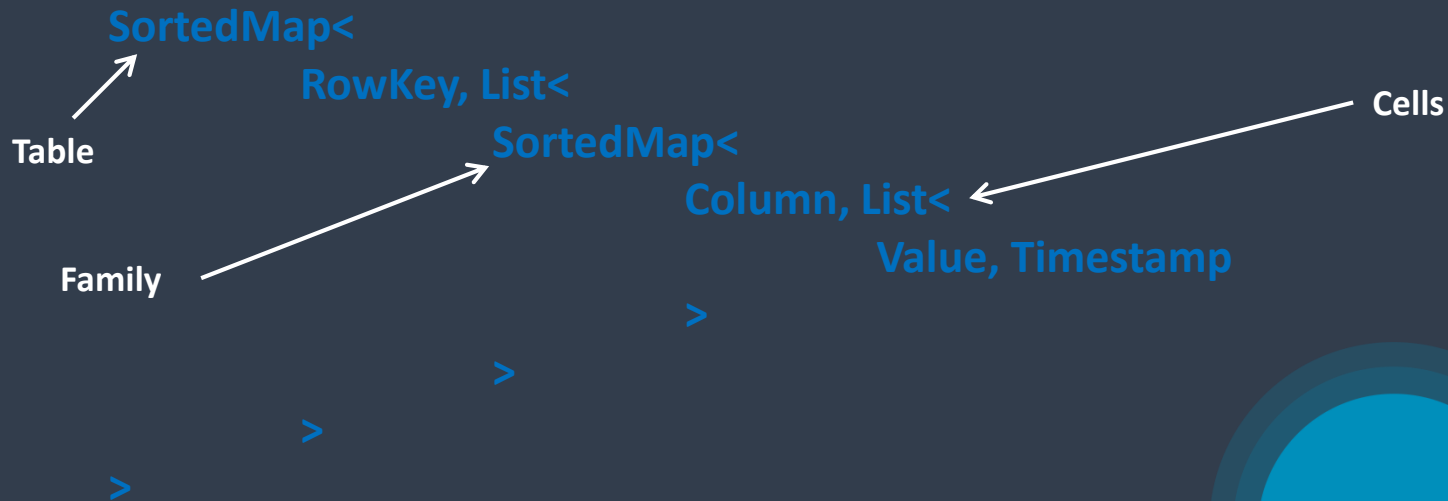
# HBASE Timestamp?

- Cells' values are versioned
  - ✓ For each cell multiple versions are kept
- 3 by default
  - ✓ Another dimension to identify your data
  - ✓ Either explicitly **timestamped** by region server or provided by the client
- Versions are stored in decreasing timestamp order
- Read the latest first – optimization to read the current value
- You can specify how many versions are kept
  - ✓ More on this later....

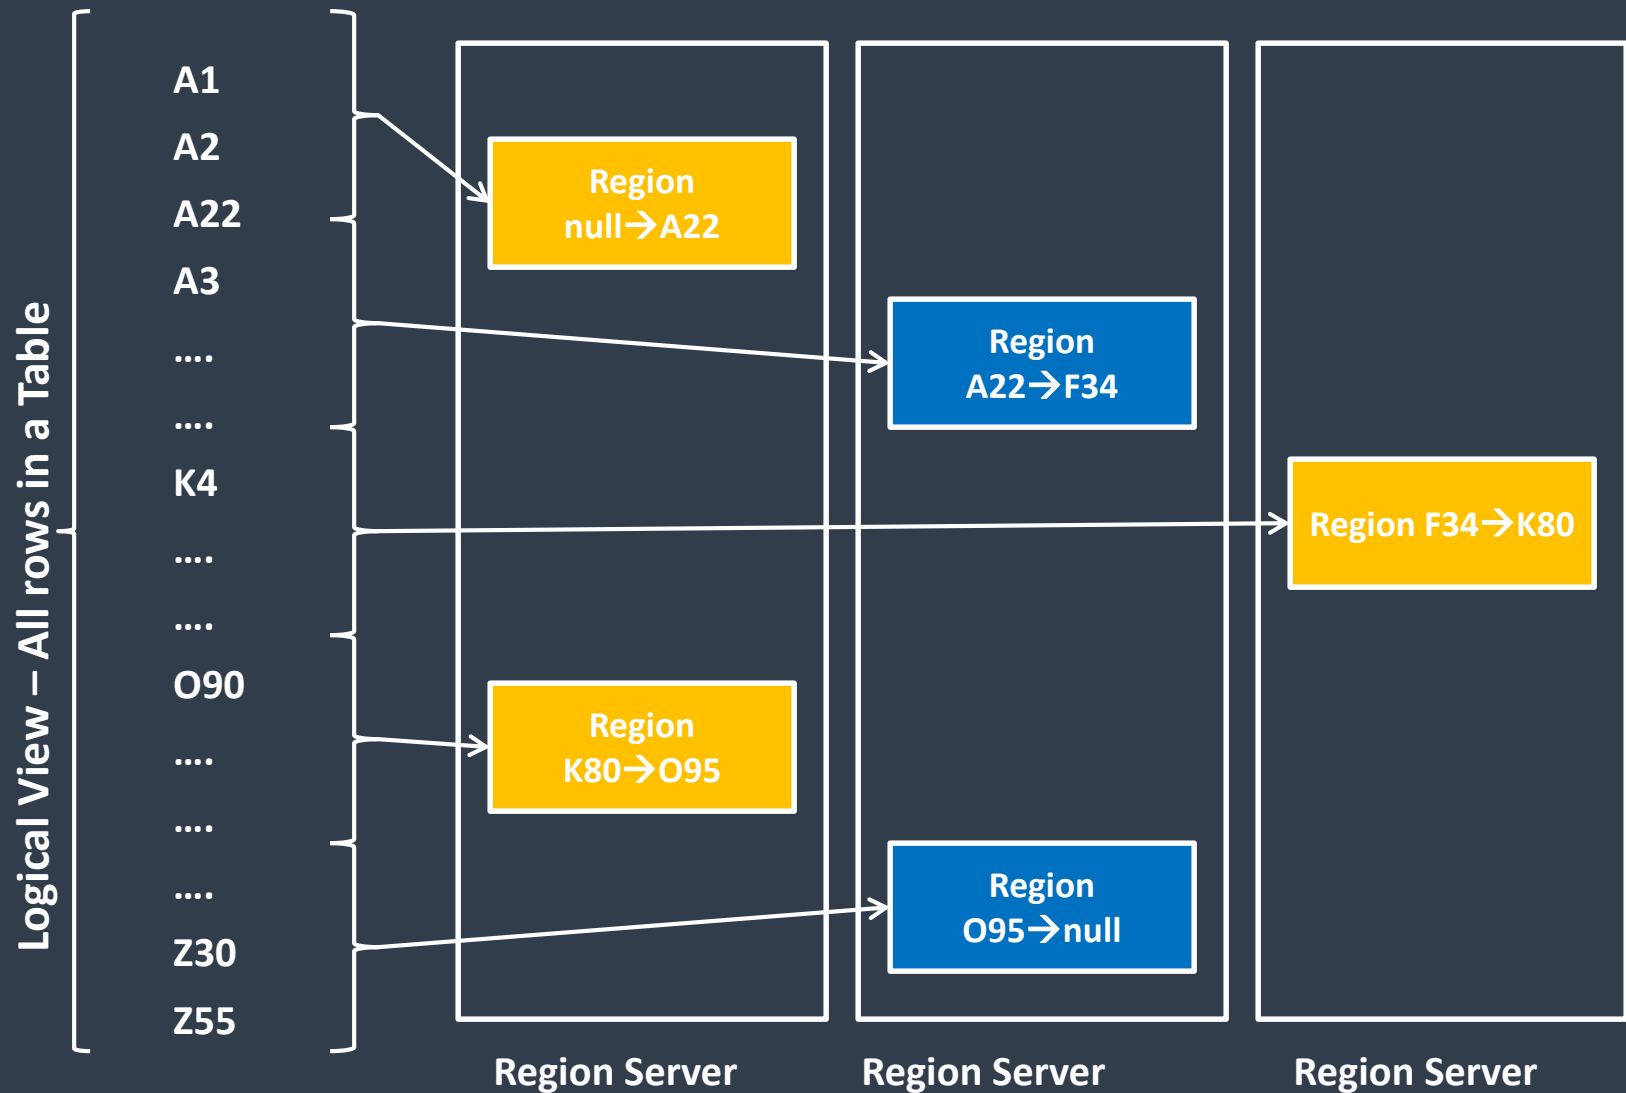


# HBASE Cells

- Value = Table+RowKey+Family+Column+Timestamp
- Programming language style:



# Row Distribution Between Region Server



# HBASE Cells

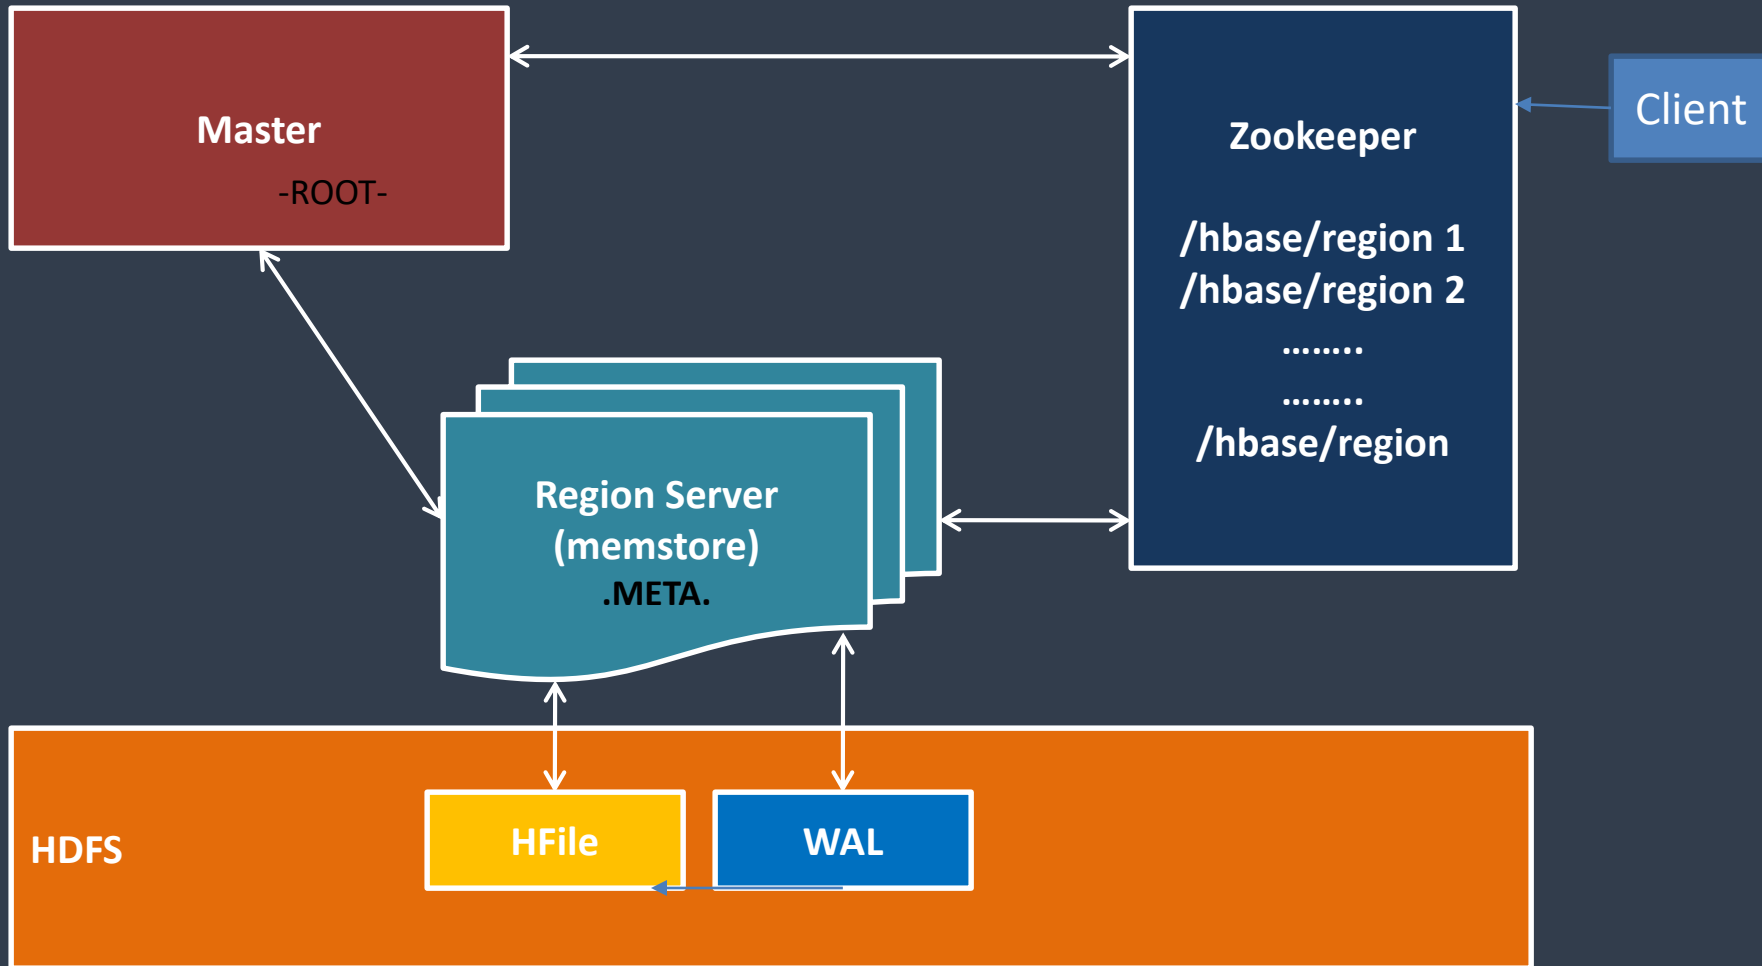
Row Key	Time Stamp	Name Family		Address Family	
		First_name	Last_name	Number	Address
row1	t1	<u>Bob</u>	<u>Smith</u>		
	t5			10	First Lane
	t10			30	Other Lane
	t15			<u>7</u>	<u>Last Street</u>
row2	t20	<u>Mary</u>	Tompson		
	t22			77	One Street
	t30		<u>Thompson</u>		

# HBASE Architecture

- Table is made of **regions**
- **Region** – a range of rows stored together
  - ✓ Single shard, used for scaling
  - ✓ Dynamically split as they become too big and merged if too small
- **Region Server**- serves one or more regions
- **Master Server** – daemon responsible for managing HBase cluster, aka Region Servers
- HBase stores its data into HDFS
  - ✓ relies on HDFS's high availability and fault-tolerance features
- HBase utilizes Zookeeper for distributed coordination




# HBASE Components



# HBASE Regions

- Region is a range of keys
  - ✓ start key → stop key (ex. **k3cod** → **odiekd**)
  - ✓ start key inclusive and stop key exclusive
- Addition of data
  - ✓ At first there is only 1 region
  - ✓ Addition of data will eventually exceed the configured maximum then region will be split
    - Default is **256MB**
  - ✓ The region is split into 2 regions at the middle key
- Regions per server depend on hardware specs, with today's hardware it's common to have:
  - ✓ 10 to 1000 regions per Region Server
  - ✓ Managing as much as 1GB to 2 GB per region

- Splitting data into regions allows
  - ✓ Fast recovery when a region fails
  - ✓ Load balancing when a server is overloaded
- May be moved between servers
  - ✓ Splitting is fast
- Reads from an original file while asynchronous process performs a split
  - ✓ All of these happen automatically without user's involvement





# HBASE Storage

- Data is stored in files called HFiles/StoreFiles
  - ✓ Usually saved in HDFS
- HFile is basically a key-value map
  - ✓ Keys are sorted lexicographically
- When data is added it's written to a log called **Write Ahead Log (WAL)** and is also stored in memory (memstore)
- Flush: when in-memory data exceeds maximum value it is flushed to an HFile
  - ✓ Data persisted to HFile can then be removed from WAL
  - ✓ Region Server continues serving read-writes during the flush operations, writing values to the WAL and memstore



- Recall that HDFS doesn't support updates to an existing file therefore HFiles are immutable

- ✓ Cannot remove key-values out of HFile(s)
- ✓ Over time more and more HFiles are created

- Delete marker is saved to indicate that a record was removed

- ✓ These markers are used to filter the data - to “hide” the deleted records
- ✓ At runtime, data is merged between the content of the HFile and WAL

- To control the number of HFiles and to keep cluster well balanced HBase periodically performs data compactions

- ✓ Minor Compaction: Smaller HFiles are merged into larger HFiles (n-way merge)

- ☐ Fast - Data is already sorted within files

- ☐ Delete markers are not applied

- ✓ Major Compaction:

- ☐ For each region merges all the files within a column-family into a single file

- ☐ Scan all the entries and apply all the deletes as necessary

# HBASE Master

- Responsible for managing regions and their locations
  - ✓ Assigns regions to region servers
  - ✓ Re-balanced to accommodate workloads
  - ✓ Recovers if a region server becomes unavailable
  - ✓ Uses Zookeeper – distributed coordination service
- Doesn't actually store or read data
  - ✓ Clients communicate directly with Region Servers
  - ✓ Usually lightly loaded
- Responsible for schema management and changes
  - ✓ Adding/Removing tables and column families



# HBASE And Zookeeper

- HBase uses Zookeeper extensively for region assignment



**“Zookeeper is a centralized service for maintaining configuration information, naming, providing distributed synchronization, and providing group services”**

**- [zookeeper.apache.org](http://zookeeper.apache.org)**

- HBase can manage Zookeeper daemons for you or you can install/manage them separately
- Learn More at <http://zookeeper.apache.org>

# HBASE And Zookeeper

- Zookeeper crash course
  - ✓ Very simple file-like API, written in Java
  - ✓ Operations on directories and files (called **Znodes**)
  - ✓ CRUD ZNodes and register for updates
    - ❑ Supports PERSISTENT and EPHERMAL Znodes
  - ✓ Clients connect with a session to Zookeeper
    - ❑ Session is maintained via heartbeat, if client fails to report then the session is expired and all the EPHERMAL nodes are deleted
    - ❑ Clients listening for updates will be notified of the deleted nodes as well as new nodes

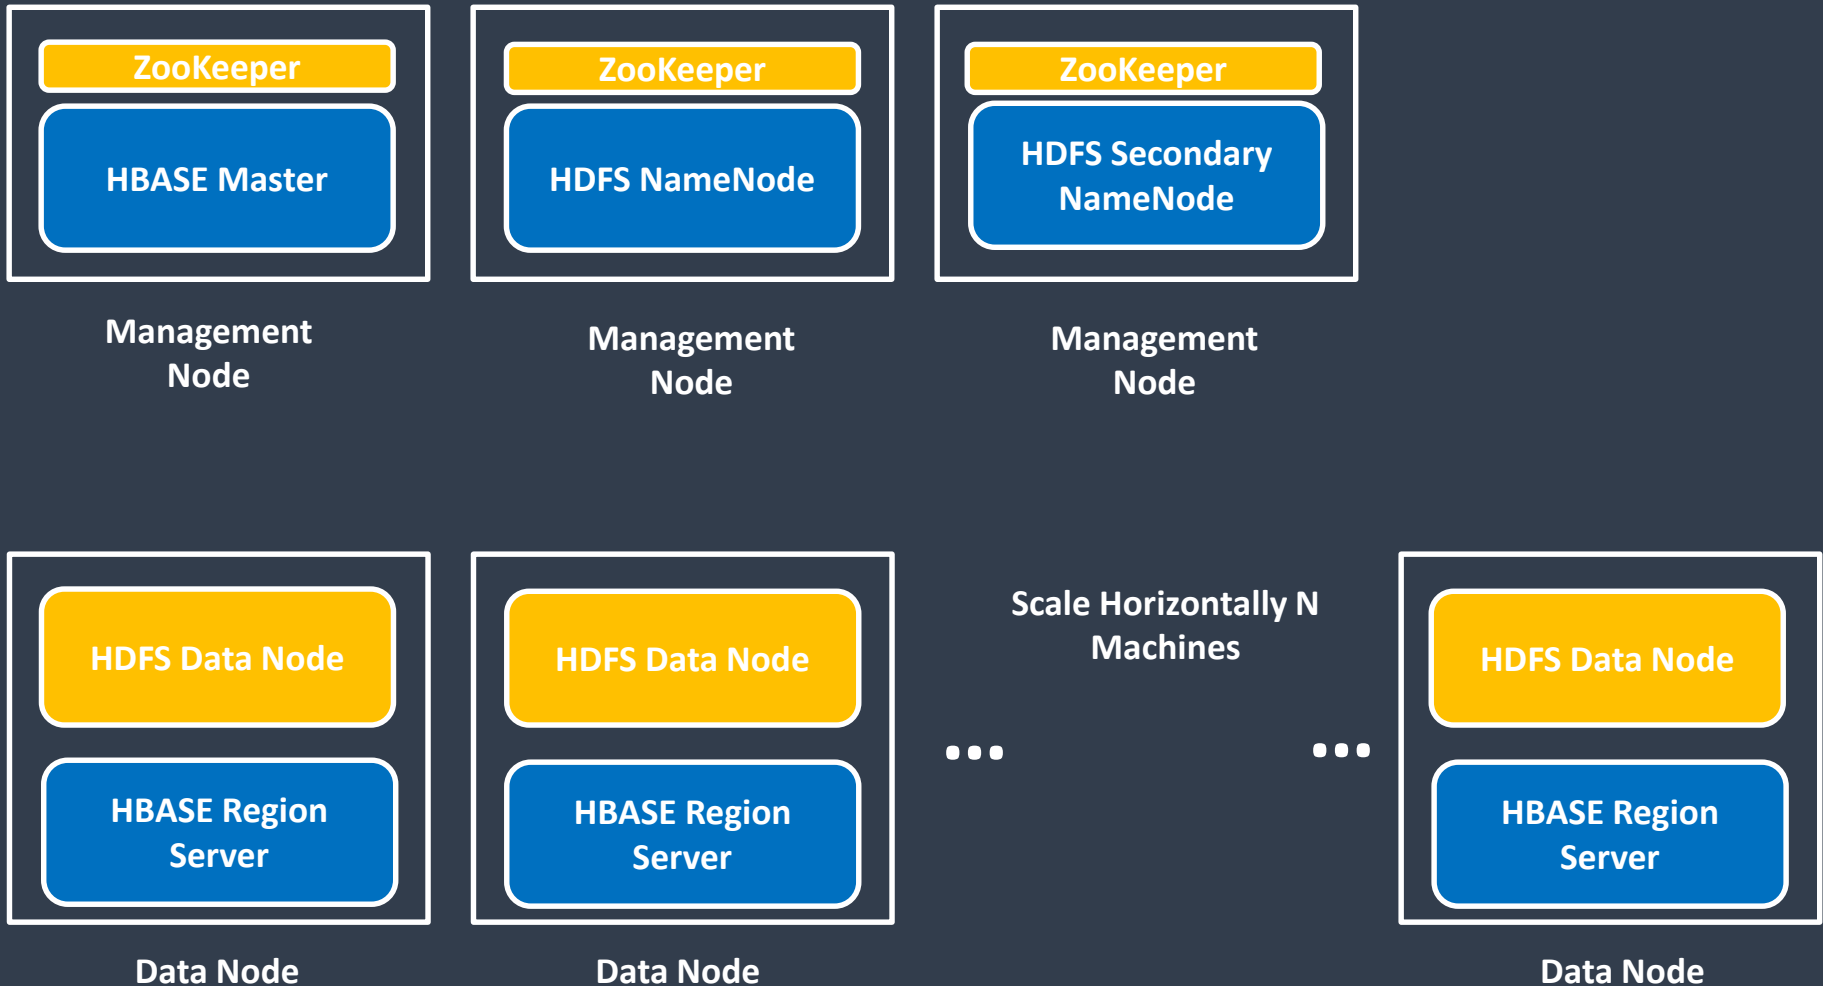


# HBASE And Zookeeper


- Each Region Server creates an ephemeral node
  - ✓ Master monitors these nodes to discover available region servers
  - ✓ Master also tracks these nodes for server failures
- Uses Zookeeper to make sure that only 1 master is registered
- HBase cannot exist without Zookeeper



# HBASE Deployment



# HBASE



# INSTALLATION



# HBASE Installation Agenda

- Learn about installation modes
- How to set-up Pseudo-Distributed Mode
- HBase Management Console
- HBase Shell
  - ✓ Define Schema
  - ✓ Create, Read, Update and Delete



# Runtime Modes

- **Local (Standalone) Mode**


- ✓ Comes Out-of-the-Box, easy to get started
- ✓ Uses local filesystem (not HDFS), NOT for production
- ✓ Runs HBase & Zookeeper in the same JVM

- **Pseudo-Distributed Mode**

- ✓ Requires HDFS
- ✓ Mimics Fully-Distributed but runs on just one host
- ✓ Good for testing, debugging and prototyping
- ✓ Not for production use or performance benchmarking!
- ✓ Development mode used in class

- **Fully-Distributed Mode**

- ✓ Run HBase on many machines
- ✓ Great for production and development clusters



# Pseudo Distributed Mode

## 1. Verify Installation Requirements

- ❖ Java, password-less SSH

## 2. Configure Java

## 3. Configure the use of HDFS

- ❖ Specify the location of Namenode
- ❖ Configure replication

## 4. Make sure HDFS is running

## 5. Start Hbase

## 6. Verify HBase is running

## 6. Verify HBASE Running

```
$ hbase shell
```

HBase Shell; enter 'help<RETURN>' for list of supported commands.

Type "exit<RETURN>" to leave the HBase Shell


Version 0.90.4-cdh3u2, r, Thu Oct 13 20:32:26 PDT 2011

```
hbase(main):001:0> list
```

TABLE

0 row(s) in 0.4070 seconds

Run a command to verify  
that cluster is actually  
running



```
$ hadoop fs -ls /hbase
```

Found 5 items


drwxr-xr-x	-	hadoop	supergroup	0	2011-12-31 13:18	/hbase/-ROOT-
drwxr-xr-x	-	hadoop	supergroup	0	2011-12-31 13:18	/hbase/.META.
drwxr-xr-x	-	hadoop	supergroup	0	2011-12-31 13:18	/hbase/.logs
drwxr-xr-x	-	hadoop	supergroup	0	2011-12-31 13:18	/hbase/.oldlogs
-rw-r--r--	1	hadoop	supergroup	3	2011-12-31 13:18	/hbase/hbase.version

HBase data and metadata is  
stored in HDFS



# HBASE Management Console

- HBase comes with web based management
  - ✓ <http://localhost:60010>
- Both Master and Region servers run web server
  - ✓ Browsing Master will lead you to region servers
    - ❑ Regions run on port 60030
- Firewall considerations
  - ✓ Opening <master\_host>:60010 in firewall is not enough
  - ✓ Have to open up <region(s)\_host>:60030 on every slave host
  - ✓ An easy option is to open a browser behind the firewall
    - ❑ SSH tunneling and **Virtual Network Computing (VNC)**



# HBASE Shells

- JRuby IRB (Interactive Ruby Shell)
  - ✓ HBase commands added
  - ✓ If you can do it in IRB you can do it in HBase shell
- [http://en.wikipedia.org/wiki/Interactive\\_Ruby\\_Shell](http://en.wikipedia.org/wiki/Interactive_Ruby_Shell)
- To run simply

```
$ <hbase_install>/bin/hbase shell
HBase Shell; enter 'help<RETURN>' for list of supported commands.
Type "exit<RETURN>" to leave the HBase Shell
Version 0.90.4-cdh3u2, r, Thu Oct 13 20:32:26 PDT 2011
hbase(main):001:0>
```

  - ✓ Puts you into IRB
  - ✓ Type 'help' to get a listing of commands
    - ❑ \$ help “command” (quotes are required)
    - > help “get”



## If Daemons are not running or Hbase is not working

- **sudo service hbase-master restart**
- **sudo service hbase-regionserver restart**
- **sudo service zookeeper-server restart**

- Quote all names
  - ✓ Table and column names
  - ✓ Single quotes for text
  - hbase> get 't1', 'myRowId'
  - ✓ Double quotes for binary
    - ❑ Use hexadecimal representation of that binary value
    - ❑ hbase> get 't1', "key\x03\x3f\xcd"
- Uses ruby hashes to specify parameters
  - ✓ {'key1' => 'value1', 'key2' => 'value2', ...}
  - ✓ Example:
 

```
hbase> get 'UserTable', 'userid1', {COLUMN => 'address:str'}
```
- HBase Shell supports various commands
  - ✓ General
    - ❑ status, version
  - ✓ Data Definition Language (DDL)
    - ❑ alter, create, describe, disable, drop, enable, exists, is\_disabled, is\_enabled, list
  - ✓ Data Manipulation Language (DML)
    - ❑ count, delete, deleteall, get, get\_counter, incr, put, scan, truncate
  - ✓ Cluster administration
    - ❑ balancer, close\_region, compact, flush, major\_compact, move, split, unassign, zk\_dump, add\_peer, disable\_peer, enable\_peer, remove\_peer, start\_replication, stop\_replication
- Learn more about each command
  - ✓ hbase> help "<command>"

# HBASE Shell — Check Status

- Display cluster's status via status command
  - ✓ `hbase> status`
  - ✓ `hbase> status 'detailed'`
- Similar information can be found on Hbase Web Management Console
  - ✓ <http://localhost:60010>

```
hbase> status
```

```
1 servers, 0 dead, 3.0000 average load
```

```
hbase> status 'detailed'
```

```
version 0.90.4-cdh3u2
```

```
0 regionsInTransition
```

```
1 live servers
```

```
    hadoop-laptop:39679 1326056194009
```

```
        requests=0, regions=3, usedHeap=30, maxHeap=998
```

```
        .META.,,1
```

```
            stores=1, storefiles=0, storefileSizeMB=0, ...
```

```
        -ROOT-,0
```

```
            stores=1, storefiles=1, storefileSizeMB=0, ...
```

```
    Blog,,1326059842133.c1b865dd916b64a6228ecb4f743 ...
```

```
0 dead servers
```



# HBASE Shell **DDL** and **DML**

Let's walk through an example

1. Create a table
  - Define column families
2. Populate table with data records
  - Multiple records
3. Access data
  - Count, get and scan
4. Edit data
5. Delete records
6. Drop table



# HBASE Shell **NameSpace**

```
create_namespace 'ns1'  
describe_namespace 'ns1'  
list_namespace  
list_namespace_tables 'default'
```

```
alter_namespace 'ns1', {METHOD => 'set', 'PROPERTY_NAME' =>  
'PROPERTY_VALUE'}
```

```
alter_namespace 'ns1', {METHOD => 'unset',  
NAME=>'PROPERTY_NAME'}
```



# 1.Create Table

- Create table called 'Blog' with the following schema
  - ✓ 2 families
    - ❑ 'info' with 3 columns: 'title', 'author', and 'date'
    - ❑ 'content' with 1 column family: 'post'

Blog		
Family:	Info:	Columns: title, author, date
	Content:	Columns:Post


# 1.Create Table

- Various options to create tables and families

- ✓ `hbase> create 't1', {NAME => 'f1', VERSIONS => 5}`
- ✓ `hbase> create 't1', {NAME => 'f1', VERSIONS => 1,TTL => 2592000, BLOCKCACHE => true}`
- ✓ `hbase> create 't1', {NAME => 'f1'}, {NAME => 'f2'},{NAME => 'f3'}`
- ✓ `hbase> create 't1', 'f1', 'f2', 'f3'`

`hbase> create 'Blog', {NAME=>'info'}, {NAME=>'content'}`

0 row(s) in 1.3580 seconds



## 2. Populate Table With Data Records

- Populate data with multiple records

Row Id	Info:title	Info:author	Info:date	Content:post
Matt-001	Elephant	Matt	2009.05.06	Do Elephants like monkeys?
Matt-002	Monkey	Matt	2011.02.14	Do monkeys like elephants?
Bob-003	Dog	Bob	1995.10.20	People Own Dogs!
Michelle-004	Cat	Michelle	1990.07.06	I have a cat!
John-005	Mouse	John	2012.01.15	Mickey Mouse

- Put command format:

```
hbase> put 'table', 'row_id', 'family:column', 'value'
```

### # insert row 1

```
put 'Blog', 'Matt-001', 'info:title', 'Elephant'  
put 'Blog', 'Matt-001', 'info:author', 'Matt'  
put 'Blog', 'Matt-001', 'info:date', '2009.05.06'  
put 'Blog', 'Matt-001', 'content:post', 'Do elephants like monkeys?'
```

...

...

### ... # insert rows 2-4

...

...

### # row 5

```
put 'Blog', 'John-005', 'info:title', 'Mouse'  
put 'Blog', 'John-005', 'info:author', 'John'  
put 'Blog', 'John-005', 'info:date', '1990.07.06'  
put 'Blog', 'John-005', 'content:post', 'Mickey mouse.'
```

Put statement per cell



# Access Data - Count

- Access Data
  - ✓ count: display the total number of records
  - ✓ get: retrieve a single row
  - ✓ scan: retrieve a range of rows
- Count is simple
  - ✓ `hbase> count 'table_name'`
  - ✓ Will scan the entire table! May be slow for a large table
    - ❑ Alternatively can run a MapReduce job (more on this later...)
      - `$ yarn jar hbase.jar rowcount`
  - ✓ Specify count to display every n rows. Default is 1000
    - ❑ `hbase> count 't1', INTERVAL => 10`



```
hbase> count 'Blog', {INTERVAL=>2}
```

```
Current count: 2, row: John-005
```

```
Current count: 4, row: Matt-002
```

```
5 row(s) in 0.0220 seconds
```

```
hbase> count 'Blog', {INTERVAL=>1}
```

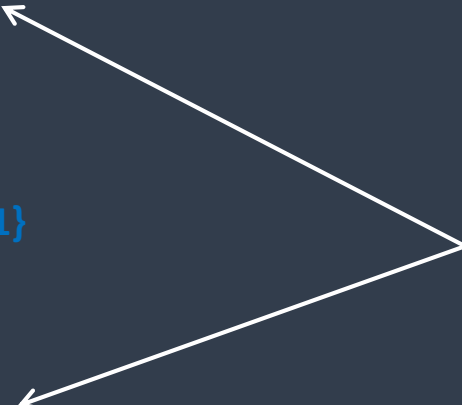
```
Current count: 1, row: Bob-003
```

```
Current count: 2, row: John-005
```

```
Current count: 3, row: Matt-001
```

```
Current count: 4, row: Matt-002
```

```
Current count: 5, row: Michelle-004
```



Affect how often  
count is displayed



# Access Data - Get

- Select single row with 'get' command
  - ✓ `hbase> get 'table', 'row_id'`
    - Returns an entire row
  - ✓ Requires table name and row id
  - ✓ Optional: timestamp or time-range, and versions
- Select specific columns
  - ✓ `hbase> get 't1', 'r1', {COLUMN => 'c1'}`
  - ✓ `hbase> get 't1', 'r1', {COLUMN => ['c1', 'c2', 'c3']}`
- Select specific timestamp or time-range
  - ✓ `hbase> get 't1', 'r1', {TIMERANGE => [ts1, ts2]}`
  - ✓ `hbase> get 't1', 'r1', {COLUMN => 'c1', TIMESTAMP => ts1}`
- Select more than one version
  - ✓ `hbase> get 't1', 'r1', {VERSIONS => 4}`



# Access Data - Get

```
hbase> get 'Blog', 'unknownRowId'
```

COLUMN

CELL

0 row(s) in 0.0250 seconds

Row Id Doesn't exist



```
hbase> get 'Blog', 'Michelle-004'
```

COLUMN

CELL

content:post

timestamp=1326061625690, value=I have a cat!

info:author

timestamp=1326061625630, value=Michelle

info:date


timestamp=1326061625653, value=1990.07.06

info:title

timestamp=1326061625608, value=Cat

4 row(s) in 0.0420 seconds

Returns ALL Columns, display 1 column per row!!!



# Access Data - Get

Narrow down to just two columns

```
hbase> get 'Blog', 'Michelle-004',  
        {COLUMN=>['info:author','content:post']}
```

COLUMN	CELL
content:post	timestamp=1326061625690, value=I have a cat!
info:author	timestamp=1326061625630, value=Michelle

2 row(s) in 0.0100 seconds

Narrow down to via columns and timestamp

```
hbase> get 'Blog', 'Michelle-004',  
        {COLUMN=>['info:author','content:post'],TIMESTAMP=>1326061625690}
```

COLUMN	CELL
content:post timestamp=1326061625690	value=I have a cat!

1 row(s) in 0.0140 seconds

Only One timestamp matches

# Access Data - Get

```
hbase> get 'Blog', 'Michelle-004',  
        {COLUMN=>'info:date', VERSIONS=>2}
```

COLUMN	CELL
info:date	timestamp=1326071670471, value=1990.07.08
info:date	timestamp=1326071670442, value=1990.07.07

2 row(s) in 0.0300 seconds

Ask for the latest two versions



```
hbase> get 'Blog', 'Michelle-004',  
        {COLUMN=>'info:date'}
```

COLUMN	CELL
info:date	timestamp=1326071670471, value=1990.07.08

1 row(s) in 0.0190 seconds

By Default only the latest version is returned



# Access Data - Scan

- Scan entire table or a portion of it
- Load entire row or explicitly retrieve column families, columns or specific cells
- To scan an entire table
  - ✓ `hbase> scan 'table_name'`
- Limit the number of results
  - ✓ `hbase> scan 'table_name', {LIMIT=>1}`
- Scan a range
  - ✓ `hbase> scan 'Blog', {STARTROW=>'startRow', STOPROW=>'stopRow'}`
  - ✓ Start row is inclusive, stop row is exclusive
  - ✓ Can provide just start row or just stop row



# Access Data - Scan

- Limit what columns are retrieved
  - ✓ `hbase> scan 'table', {COLUMNS=>['col1', 'col2']}`
- Scan a time range
  - ✓ `hbase> scan 'table', {TIMERANGE => [1303, 13036]}`
- Limit results with a filter
  - ✓ `hbase> scan 'Blog', {FILTER =>`  
`org.apache.hadoop.hbase.filter.ColumnPaginationFilter.new(1, 0)}`
  - ✓ More about filters later
  - ✓ `scan 'blog2',{FILTER => "ValueFilter(=,'binaryprefix:Mickey')"`
  - ✓ `scan 'blog2',{FILTER => "FirstKeyOnlyFilter()"`
  - ✓ `scan 'blog2',{FILTER => "KeyOnlyFilter()"`
  - ✓ `scan 'blog2',{FILTER => "(PrefixFilter('john'))"`

<http://hadooptutorial.info/hbase-functions-cheat-sheet/>

Scan the entire table, grab ALL the columns


**hbase(main):014:0> scan 'Blog'**

**ROW COLUMN+CELL**

Bob-003 column=content:post,	timestamp=1326061625569, value=People own dogs!
Bob-003 column=info:author,	timestamp=1326061625518, value=Bob
Bob-003 column=info:date,	timestamp=1326061625546, value=1995.10.20
Bob-003 column=info:title,	timestamp=1326061625499, value=Dog
John-005 column=content:post,	timestamp=1326061625820, value=Mickey mouse.
John-005 column=info:author,	timestamp=1326061625758, value=John
...	
Michelle-004 column=info:author,	timestamp=1326061625630, value=Michelle
Michelle-004 column=info:date,	timestamp=1326071670471, value=1990.07.08
Michelle-004 column=info:title,	timestamp=1326061625608, value=Cat

5 row(s) in 0.0670 seconds

Stop row is exclusive, row ids that start with John will not be included



```
hbase> scan 'Blog', {STOPROW=>'John'}
```

ROW	COLUMN+CELL	
-----	-------------	--

Bob-003	column=content:post,	timestamp=1326061625569, value=People own dogs!
Bob-003	column=info:author,	timestamp=1326061625518, value=Bob
Bob-003	column=info:date,	timestamp=1326061625546, value=1995.10.20
Bob-003	column=info:title,	timestamp=1326061625499, value=Dog

1 row(s) in 0.0410 seconds



Only retrieve 'info:title' column



```
hbase> scan 'Blog', {COLUMNS=>'info:title',  
                      STARTROW=>'John', STOPROW=>'Michelle'}
```

ROW	COLUMN+CELL	
John-005	column=info:title,	timestamp=1326061625728, value=Mouse
Matt-001	column=info:title,	timestamp=1326061625214, value=Elephant
Matt-002	column=info:title,	timestamp=1326061625383, value=Monkey

3 row(s) in 0.0290 seconds

# Edit Data

- Put command inserts a new value if row id doesn't exist
- Put updates the value if the row does exist
- But does it really update?
  - ✓ Inserts a new version for the cell
  - ✓ Only the latest version is selected by default
  - ✓ N versions are kept per cell
    - ❑ configured per family at creation:
      - `hbase> create 'table', {NAME => 'family', VERSIONS => 7}`
    - ❑ 3 versions are kept by default



```
hbase> put 'Blog', 'Michelle-004', 'info:date', '1990.07.06'
```


```
0 row(s) in 0.0520 seconds
```

```
hbase> put 'Blog', 'Michelle-004', 'info:date', '1990.07.07'
```

```
0 row(s) in 0.0080 seconds
```

```
hbase> put 'Blog', 'Michelle-004', 'info:date', '1990.07.08'
```

```
0 row(s) in 0.0060 seconds
```



Update the same exact row with a different value

```
hbase> get 'Blog', 'Michelle-004',
```

```
{COLUMN=>'info:date', VERSIONS=>3}
```

COLUMN	CELL
info:date	timestamp=1326071670471, value=1990.07.08
info:date	timestamp=1326071670442, value=1990.07.07
info:date	timestamp=1326071670382, value=1990.07.06

```
3 row(s) in 0.0170 seconds
```



Keeps three versions of each cell by default

```
hbase> get 'Blog', 'Michelle-004',
```

```
{COLUMN=>'info:date', VERSIONS=>2}
```

COLUMN	CELL
info:date	timestamp=1326071670471, value=1990.07.08
info:date	timestamp=1326071670442, value=1990.07.07

2 row(s) in 0.0300 seconds



Asks for the latest two versions

```
hbase> get 'Blog', 'Michelle-004',
```

```
{COLUMN=>'info:date'}
```

COLUMN	CELL
info:date	timestamp=1326071670471, value=1990.07.08

1 row(s) in 0.0190 seconds



By default only the latest version is returned

# Delete Records

- Delete cell by providing table, row id and column coordinates
  - ✓ delete 'table', 'rowId', 'column'
  - ✓ Deletes all the versions of that cell
  - ✓ **deleteall** command for deleting the entire row
  - ✓ **truncate** command for deleting all the rows at a time but keeping the schema
- Optionally add timestamp to only delete versions before the provided timestamp
  - ✓ **Delete 'table', 'rowId', 'column', timestamp**

✓ scan 'blog',{RAW => true}

Matt-001


type>DeleteColumn

Matt-001

value=Ramesh

column=info:author, timestamp=1467774783983

column=info:author, timestamp=1467772792122



# Delete Records

```
hbase> get 'Blog', 'Bob-003', 'info:date'
```

```
COLUMN
```

```
CELL
```

```
info:date          timestamp=1326061625546, value=1995.10.20
```

```
1 row(s) in 0.0200 seconds
```

```
hbase> delete 'Blog', 'Bob-003', 'info:date'
```

```
0 row(s) in 0.0180 seconds
```

```
hbase> get 'Blog', 'Bob-003', 'info:date'
```

```
COLUMN
```

```
CELL
```

```
0 row(s) in 0.0170 seconds
```

# Delete Records

```
hbase> get 'Blog', 'Michelle-004',
```

```
{COLUMN=>'info:date', VERSIONS=>3}
```

COLUMN	CELL
info:date	timestamp=1326254742846, value=1990.07.08
info:date	timestamp=1326254739790, value=1990.07.07
info:date	timestamp=1326254736564, value=1990.07.06

3 row(s) in 0.0120 seconds



3 versions

```
hbase> delete 'Blog', 'Michelle-004', 'info:date', 1326254739791
```

0 row(s) in 0.0150 seconds



1 millisecond after the  
second version

```
hbase> get 'Blog', 'Michelle-004',
```

```
{COLUMN=>'info:date', VERSIONS=>3}
```

COLUMN	CELL
info:date	timestamp=1326254742846, value=1990.07.08

1 row(s) in 0.0090 seconds



After the timestamp provided  
at delete statement

# Drop Table

- Must disable before dropping
  - ✓ puts the table “offline” so schema based operations can be performed
  - ✓ hbase> disable 'table\_name'
  - ✓ hbase> drop 'table\_name'
- For a large table it may take a long time....

```
hbase> list
```

```
TABLE
```

```
Blog
```

```
1 row(s) in 0.0120 seconds
```

Take the table offline for  
schema modifications



```
hbase> disable 'Blog'
```

```
0 row(s) in 2.0510 seconds
```

```
hbase> drop 'Blog'
```

```
0 row(s) in 0.0940 seconds
```

```
hbase> list
```

```
TABLE
```

```
0 row(s) in 0.0200 seconds
```



# Bulk Import

```
[cloudera@quickstart ~]$ hadoop fs -appendToFile - /user/cloudera/hbase_bulkimport_tsv  
123      emp1      sse  
345      emp2      tl  
456      emp3      sse  
hbase(main):002:0> create 'emp','info','desc'  
[cloudera@quickstart ~]$ hbase org.apache.hadoop.hbase.mapreduce.ImportTsv -  
Dimporttsv.columns=HBASE_ROW_KEY,info:name,desc:role emp  
/user/cloudera/hbase_bulkimport_tsv
```

Onprem:

```
hbase org.apache.hadoop.hbase.mapreduce.Export ns1:FACT_DETAIL  
/user/nreuser/export/fact_detail/fact_detail
```

copy the file to cloud edge node

Cloud

```
hbase org.apache.hadoop.hbase.mapreduce.Import ns1:FACT_DETAIL  
/user/selvka12/folder/fact_detail/
```

<http://hadooptutorial.info/forums/topic/hbase-bulk-loading-with-importtsv/>

# Accessing Hbase table in Hive

```
hive> CREATE EXTERNAL TABLE hbase_table_emp(key string, name string, role string)
> STORED BY 'org.apache.hadoop.hive.hbase.HBaseStorageHandler'
> WITH SERDEPROPERTIES ("hbase.columns.mapping" = "info:name,desc:role")
> TBLPROPERTIES("hbase.table.name" = "emp");
```

```
hive> select * from hbase_table_emp;
```

123	emp1	sse
345	emp2	tl
456	emp3	sse

Load data into hive external table which will be inserted into Hbase tables as well

All these tables will be visible and queried from Impala shell as well



# Agenda

- Create via Put method
- Read via Get method
- Update via Put method
- Delete via Delete method



# JAVA Client API Overview

- **HBase** is written in Java
  - ✓ No surprise that it has a Java Native API
- Supports programmatic access to Data Manipulation Language (DML)
  - ✓ CRUD operations plus more
- Everything that you can do with **HBase** Shell and more....
- Java Native API is the fastest way to access **HBase**

## Using Client API

1. Create a Configuration object
  - ✓ Recall Configuration from HDFS object
  - ✓ Adds HBase specific props
2. Construct HTable
  - ✓ Provide Configuration object
  - ✓ Provide table name
3. Perform operations
  - ✓ Such as put, get, scan, delete, etc...
4. Close HTable instance
  - ✓ Flushes all the internal buffers
  - ✓ Releases all the resources



**1. Create a Configuration object**

```
Configuration conf = HbaseConfiguration.create();
```

**2. Construct HTable**

```
HTable hTable = new HTable(conf, tableName);
```

**3. Perform operations**

```
hTable.getTableNames();
```

**4. Close HTable instance**

```
hTable.close();
```



# ConstructHTable.java

```
public class ConstructHTable
{
    public static void main(String[] args) throws IOException
    {
        Configuration conf = HbaseConfiguration.create();
```

Seeds configuration object with required information to establish client connection

```
        HTable hTable = new HTable(conf, "-ROOT-");
```

Table name

```
        System.out.println("Table is: " +
            Bytes.toString(hTable.getTableNames()));
```

```
        hTable.close();
    }
}
```

Release all the resource

# ConstructHTable.java Output

```
$ yarn jar $PLAY_AREA/HadoopSamples.jar hbase.ConstructHTable
```

```
12/01/15 13:22:03 INFO zookeeper.ZooKeeper: Client  
environment:zookeeper.version=3.3.3-cdh3u2--1, built on  
10/14/2011 03:25 GMT
```

```
...
```

```
...
```

```
...
```

```
12/01/15 13:22:03 INFO zookeeper.ClientCnxn: Session  
establishment complete on server localhost/127.0.0.1:2181,  
sessionid = 0x134e27760560013, negotiated timeout = 40000
```

```
Table is: -ROOT-
```



# 1.Create Configuration Object

- Client Code Configuration
- HbaseConfiguration extends Hadoop's Configuration class
  - ✓ Still fully compatible with Configuration
- How did HbaseConfiguration.create() seed Configuration object?
  - ✓ Loads hbase-default.xml and hbase-site.xml from Java CLASSPATH
    - ☐ hbase-default.xml is packaged inside HBase jar
    - ☐ hbase-site.xml will need to be added to the CLASSPATH
    - ☐ hbase-site.xml overrides properties in hbase-default.xml

- How did hbase-site.xml get on CLASSPATH?

- ✓ Recall that we executed the code via yarn script

```
$ yarn jar HadoopSamples.jar hbase.ConstructHTable
```

- ✓ Hadoop's scripts are configured to put hbase's CLASSPATH onto it's CLASSPATH
- ✓ Specified in <hadoop\_install>/conf/hadoop-env.sh

```
export HADOOP_CLASSPATH=  
    $HBASE_HOME/*:$HBASE_HOME/conf:$HADOOP_CLASSPATH
```

- ✓ To check what's on Hadoop's CLASSPATH
  - ❑ \$ yarn classpath
  - ❑ \$ yarn classpath | grep hbase

- Creating **HTable** instance is not free
  - ✓ Actually quite costly – scans catalog .META. Table
    - Checks that table exists and enabled
  - ✓ Create once (per thread) and re-use for as long as possible
  - ✓ If you find yourself constructing many instances consider using HTablePool (utility to re-use multiple Htable instances)
- **HTable** is NOT thread safe
  - ✓ Create 1 instance per thread
- **HTable** supports CRUD batch operations
  - ✓ Not atomic
  - ✓ For performance and convenience

# Create / Save Data To HBase

1. Construct HTable instance
  - ✓ Create Put instance
2. Add cell values and their coordinates
  - ✓ Specify family:column as a coordinate
3. Call put on HTable instance
4. Close HTable



# 1. Construct HTable

- Create Configuration
- Construct HTable

```
Configuration conf = HBaseConfiguration.create();  
HTable hTable = new HTable(conf, "HBaseSamples");
```



## 2. Create Put Instance

- Put is a save operation for a single row
- Must provide a row id to the constructor
  - ✓ Row id is raw bytes: can be anything like number or UUID
    - ❑ You are responsible for converting the id to bytes
    - ❑ HBase comes with a helper class Bytes that provides static methods which handles various conversions from and to bytes
      - `org.apache.hadoop.hbase.util.Bytes`

```
Put put1 = new Put(Bytes.toBytes("row1"));
```

- ✓ Optionally can provide cell's timestamp and an instance of RowLock

```
Put put2 = new Put(Bytes.toBytes("row2"), timestamp);  
Put put3 = new Put(Bytes.toBytes("row3"), rowLock);  
Put put4 = new Put(Bytes.toBytes("row4"), timestamp, rowLock);
```

### 3.Add Cell Value And Their Coordinates

- Add columns to save to Put instance
  - ✓ Provide family:value coordinate and optional timestamp
  - ✓ Few options of the add methods
    - ❑ `Put.add(family, column, value)`
    - ❑ `Put.add(family, column, timestamp, value)`
    - ❑ `Put.add(KeyValue kv)`
  - ✓ Family, column, and value are raw binary
  - ✓ Client's responsibility to convert to binary format
  - ✓ KeyValue class as its internal cell's representation
    - ❑ For advanced usage, not usually required

```
put1.add(toBytes("test"), toBytes("col1"), toBytes("val1"));  
put1.add(toBytes("test"), toBytes("col2"), toBytes("val2"));
```

## 4.Call Put On HTable Instance

- Provide initialized Put object to HTable
- The operation is synchronous

...  
`hTable.put(put1);`





## 5. Close HTable

- Release resource held by Htable
- Inform HConnectionManager that this instance won't be using connection  
`hTable.close();`

- Utilize try/finally block

```
HTable hTable = new HTable(conf, "HBaseSamples");  
try {  
    // to stuff with table  
} finally {  
    hTable.close();  
}
```

- Most examples emit try/finally constructs in favor of readability

# PutExample.java

Static import of Bytes class

```
import static org.apache.hadoop.hbase.util.Bytes.*;
```

```
public class PutExample
```

```
{
```

```
    public static void main(String[] args) throws IOException
```

```
    {
```

```
        Configuration conf = HBaseConfiguration.create();
```

```
        HTable hTable = new HTable(conf, "HBaseSamples");
```

```
        Put put1 = new Put(toBytes("row1"));
```

← Create put with id "row1"

```
        put1.add(toBytes("test"), toBytes("col1"), toBytes("val1"));
```

```
        put1.add(toBytes("test"), toBytes("col2"), toBytes("val2"));
```

```
        hTable.put(put1);
```

← Add "val1" to test:col1 column  
Add "val2" to test:col2 column

```
        hTable.close();
```

← Save row to HBase

```
    }
```

# PutExample.java Output

```
$ yarn jar HadoopSamples.jar hbase.PutExample
```

```
$ hbase shell
```

```
HBase Shell; enter 'help<RETURN>' for list of supported commands.
```

```
Type "exit<RETURN>" to leave the HBase Shell
```

```
Version 0.90.4-cdh3u2, r, Thu Oct 13 20:32:26 PDT 2011
```

```
hbase(main):001:0> get 'HBaseSamples', 'row1'
```

```
ROW      COLUMN+CELL
```

```
row1     column=test:col1, timestamp=1326663102473, value=val1
```

```
row1     column=test:col2, timestamp=1326663102473, value=val2
```

```
1 row(s) in 0.3340 seconds
```

# Retrieving Data

- API supports
  - ✓ Get a single row by id
  - ✓ Get a set of rows by a set of row ids
    - ❑ Implemented via batching and will be covered later
  - ✓ Scan an entire table or a sub set of rows
    - ❑ To scan a portion of the table provide start and stop row ids
    - ❑ Recall that row-ids are ordered by raw byte comparison
    - ❑ In case of string based ids, the order is alphabetical
- That's it
  - ✓ Very limited simple API



# Retrieving Single Row

1. Construct HTable instance
2. Create Get instance
3. Optionally narrow down result
  - ✓ Specify family:column coordinate
  - ✓ Optionally add filters
4. Request and get results
  - ✓ Call get on HTable
  - ✓ Result instance is returned and will contain the data
5. Close HTable



## 2.Create Get Instance

- Retrieve a single row
- Construct a Get Instance by providing row id
  - ✓ Row id is in raw binary format
- Optional parameter for a row lock

```
Get get = new Get(toBytes("row1"));
```



### 3. Optionally Narrow Down The Result

- Only retrieve the data that you need
  - ✓ If not specified then an entire row is retrieved
  - ✓ Important, as HBase allows you to scale to millions of rows
  - ✓ Can narrow down by family, column(s), time range and max versions
  - ✓ Can provide more than one narrow down criteria
  - ✓ Family and column name parameters are in raw bytes
- Narrow down by family
  - ✓ `get.addFamily(family)`
- Narrow down by column
  - ✓ `get.addColumn(family, column)`



- Narrow down by time range
  - ✓ `get.setTimeRange(minStamp, maxStamp)`
- Specify number of versions returned
  - ✓ `get.setMaxVersions(maxVersions)`
  - ✓ By default set to 1: only returns the latest version
- Can retrieve multiple families and columns
  - ✓ `get.addFamily(family)`
  - ✓ `get.addFamily(family1)`
  - ✓ `get.addColumn(family2, column1)`
  - ✓ `get.addColumn(family2, column2)`
  - ✓ `get.setTimeRange(minStamp, maxStamp)`





## 4.Request And Get Results

- Utilize get methods on HTable
  - ✓ Provide assembled Get instance
  - ✓ Returns Result object with all the matching cells

```
Result result = hTable.get(get);  
byte [] rowId = result.getRow();  
byte [] val1 =  
    result.getValue(toBytes("test"), toBytes("col1"));  
byte [] val2 =  
    result.getValue(toBytes("test"), toBytes("col2"));
```

- Result class

- ✓ Allows you to access everything returned
- ✓ Result is NOT Thread safe

- Methods of interest

- ✓ **Result.getRow()** - get row's id
- ✓ **Result.getValue(family, column)** - get a value for a chosen cell
- ✓ **Result.isEmpty()** - true if the result is empty false otherwise
- ✓ **Result.size()** - returns number of cells
- ✓ **Result.containsColumn(family:column)** true if column exists
- ✓ There are a number of methods that provide access to underlying KeyValue objects
  - are for advanced usage and usually not required

# GetExample.java

```
public static void main(String[] args) throws IOException  
{
```

```
    Configuration conf = HBaseConfiguration.create();  
    HTable hTable = new HTable(conf, "HBaseSamples");
```

```
    Get get = new Get(toBytes("row1"));  
    Result result = hTable.get(get);  
    print(result);
```

Get the entire row



```
    get.addColumn(toBytes("test"), toBytes("col2"));  
    result = hTable.get(get);  
    print(result);
```

Select a single column test:col2



```
    hTable.close();
```

```
}
```

# GetExample.java

```
private static void print(Result result)
{
    System.out.println("-----");
    System.out.println("RowId: " + Bytes.toString(result.getRow()));
    byte [] val1 = result.getValue(toBytes("test"), toBytes("col1"));
    System.out.println("test1:col1="+Bytes.toString(val1));
    byte [] val2 = result.getValue(toBytes("test"), toBytes("col2"));
    System.out.println("test1:col2="+Bytes.toString(val2));
}
```

Retrieve row id

Print value test:col1 column

Print value test:col2 column

# GetExample.java Output

```
$ yarn jar $PLAY_AREA/HadoopSamples.jar hbase.GetExample
```

```
...
```

```
...
```

```
-----  
RowId: row1
```

```
test1:col1=val1
```

```
test1:col2=val2  
-----
```

```
RowId: row1
```

```
test1:col1=null
```

```
test1:col2=val2
```

test1:col1 wasn't selected the second time



# Deleting Data

- Deletes are per-row-basis
- Supports batching
  - ✓ Batching is not atomic, for performance and for convenience
  - ✓ More on that later..

1. Construct HTable instance
2. Create and Initialize Delete
3. Call delete on HTable
  - ✓ `htable.delete(delete);`
4. Close HTable

## 2.Create And Initialize Delete

- Construct a Delete instance
  - ✓ Similar to Get or Put
  - ✓ Delete(byte[] row)
    - ❑ Provide a row id to delete/modify
  - ✓ Delete(byte[] row, long timestamp, RowLock rowLock)
    - ❑ Optional timestamp and RowLock
- Optionally narrow down the Deletes

```
Delete delete1 = new Delete(toBytes("anotherRow"));  
delete1.deleteColumns(toBytes("family"), toBytes("loan"));  
delete1.deleteFamily(toBytes("family"));
```

- Narrow down what to delete for a row
  - ✓ If nothing provided then entire row is deleted
  - ✓ Delete a subset of a row by narrowing down
    - ❑ `public Delete deleteFamily(byte[] family)`
    - ❑ `public Delete deleteColumn(byte[] family, byte[] qualifier)`
    - ❑ `public Delete deleteColumns(byte[] family, byte[] qualifier)`
  - ✓ Notice `deleteColumn` **VS** `deleteColumns`
    - ❑ `deleteColumns` deletes ALL the versions of the cell but `deleteColumn` only deletes the latest
  - ✓ Most of the methods are overloaded to also take timestamp
    - ❑ Deletes everything on or before the provided timestamp
    - ❑ `deleteColumn` is an exception where only the exact timestamp match is removed



# DeleteExample.java

```
public static void main(String[] args) throws IOException  
{
```

```
    Configuration conf = HBaseConfiguration.create();  
    HTable hTable = new HTable(conf, "HBaseSamples");
```

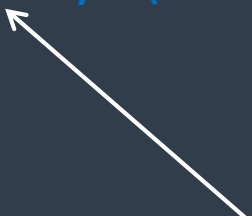
```
    Delete delete = new Delete(toBytes("rowToDelete"));  
    hTable.delete(delete);
```

Delete an entire row



```
    Delete delete1 = new Delete(toBytes("anotherRow"));  
    delete1.deleteColumns(toBytes("metrics"), toBytes("loan"));  
    hTable.delete(delete1);
```

Delete one cell rowId "anotherRow"  
and column metrics:loan



```
    hTable.close();  
}
```



## JAVA ADMIN API

# Agenda

- Create Table
- Drop Table



## JAVA Admin API

- Just like HTable is for client API HBaseAdmin is for administrative tasks
  - ✓ `org.apache.hadoop.hbase.client.HBaseAdmin`
- Recall that only Table and Family names have to be pre-defined
  - ✓ Columns can be added/deleted dynamically
  - ✓ HBase scheme roughly equals table definitions and their column families

# Create Table And Column Families

1. Construct HBaseAdmin instance
2. Create Table's schema
  - ✓ Represented by HTableDescriptor class
  - ✓ Add column families to table descriptor (HColumnDescriptor)
3. Execute create via HBaseAdmin class

## 1. Construct HBase Admin Instance

- HbaseAdmin's constructor requires an instance of Configuration object
  - ✓ Similar to HTable
  - ✓ We already know how to do that

```
Configuration conf = HBaseConfiguration.create();  
HBaseAdmin admin = new HBaseAdmin(conf);
```

## 2.Create Table Description

- `org.apache.hadoop.hbase.HTableDescriptor`
  - ✓ Serves as a container for table name and column families
  - ✓ Most importantly, add one or more column families
    - ❑ `org.apache.hadoop.hbase.HColumnDescriptor`
    - ❑ `HColumnDescriptor` serves as a container for column family name, compressions settings, number of versions, in-memory setting, and block size

```
HTableDescriptor table = new HTableDescriptor(toBytes("Table"));  
HColumnDescriptor family = new HColumnDescriptor(toBytes("f"));  
table.addFamily(family);  
HColumnDescriptor family1 = new HColumnDescriptor(toBytes("f1"));  
table.addFamily(family1);
```

## 3.Execute Create Via HBase Admin

- `HBaseAdmin` creates a table via `createTable` method
  - ✓ Synchronous operation

```
admin.createTable(table);
```

```
public static void main(String[] args) throws IOException
{
    Configuration conf = HBaseConfiguration.create();
    HBaseAdmin admin = new HBaseAdmin(conf);


    String name = "NewTable";
    byte [] tableName = toBytes(name);

    HTableDescriptor table = new HTableDescriptor(tableName);
    HColumnDescriptor family =
        new HColumnDescriptor(toBytes("new_family"));
    table.addFamily(family);

    System.out.println("Table "+name+" exist: " +
        admin.tableExists(tableName)) ;

    System.out.println("Creating "+name+" table...");
    admin.createTable(table);
    System.out.println("Table "+name+" exist: " +
        admin.tableExists(tableName)) ;
}
```

Descriptor for  
NewTable:new\_family



```
$ yarn jar $PLAY_AREA/HadoopSamples.jar hbase.CreateTableExample
```

```
...
```

```
Table NewTable exist: false
```

```
Creating NewTable table...
```

```
Table NewTable exist: true
```

```
$ hbase shell
```

```
hbase> describe 'NewTable'
```

```
DESCRIPTION
```

```
{NAME => 'NewTable', FAMILIES => [{NAME => 'new_family',  
BLOOMFILTER => 'NONE', REPLICATION_SCOPE => '0',  
COMPRESSION => 'NONE', VERSIONS => '3', TTL =>  
'2147483647', BLOCKSIZE => '65536', IN_MEMORY => 'false',  
BLOCKCACHE => 'true'}]}
```

```
ENABLED
```

```
true
```

```
1 row(s) in 0.0400 seconds
```

# Drop Table

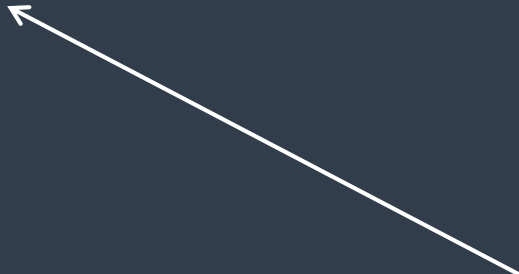
1. Construct HBaseAdmin instance
2. Disable table
  - ✓ Table must be taken offline in order to perform any schema modifications
3. Delete table

```
public static void main(String[] args) throws IOException
{
    Configuration conf = HBaseConfiguration.create();

    HBaseAdmin admin = new HBaseAdmin(conf);
    byte [] tableName = toBytes("NewTable");

    admin.disableTable(tableName);

    admin.deleteTable(tableName);
}
```



Bytes utility class is imported with 'static' keyword:  
`import static org.apache.hadoop.hbase.util.Bytes.toBytes;`



# JAVA Client Example

```
package com.test.hbase;

import java.io.IOException;

import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.hbase.HBaseConfiguration;
import org.apache.hadoop.hbase.client.Get;
import org.apache.hadoop.hbase.client.HTable;
import org.apache.hadoop.hbase.client.Put;
import org.apache.hadoop.hbase.client.Result;
import org.apache.hadoop.hbase.util.Bytes;

import org.apache.hadoop.hbase.HColumnDescriptor;
import org.apache.hadoop.hbase.HTableDescriptor;
import org.apache.hadoop.hbase.client.HBaseAdmin;
import org.apache.hadoop.hbase.client.ResultScanner;
import org.apache.hadoop.hbase.client.Scan;

public class HBaseClientExample {
```

```

/**
 * @param args
 */
@SuppressWarnings({ "deprecation", "resource" })
public static void main(String[] args) throws IOException {
    Configuration config = HBaseConfiguration.create();
    // Create table

    HBaseAdmin admin = new HBaseAdmin(config);
    HTableDescriptor htd = new HTableDescriptor("TestHtable");
    HColumnDescriptor hcd1 = new HColumnDescriptor("CF1");
    HColumnDescriptor hcd2 = new HColumnDescriptor("CF2");
    htd.addFamily(hcd1);
    htd.addFamily(hcd2);
    admin.createTable(htd);

    //List down the tables in HBase
    byte[] tablename = htd.getName();
    HTableDescriptor[] tables = admin.listTables();
    if (tables.length != 1 && Bytes.equals(tablename, tables[0].getName())){
        //throw new IOException("Failed create of table");
        System.out.println("table name is " + tables[0].getName());
    }
}

```

// Run some operations -- a put, a get, and a scan -- against the table.

```
HTable table = new HTable(config, tablename);  
byte[] row1 = Bytes.toBytes("row1");  
Put p1 = new Put(row1);  
byte[] databytes = Bytes.toBytes("CF1");  
p1.add(databytes, Bytes.toBytes("col1"), Bytes.toBytes("value1"));  
table.put(p1);  
databytes = Bytes.toBytes("CF2");  
p1.add(databytes, Bytes.toBytes("col2"), Bytes.toBytes("value2"));  
table.put(p1);
```

```
Get g = new Get(row1);  
Result result = table.get(g);  
System.out.println("Get: " + result);
```

```
Scan scan = new Scan();  
ResultScanner scanner = table.getScanner(scan);  
try {  
    for (Result scannerResult : scanner) {  
        System.out.println("Scan: " + scannerResult);  
    }  
}
```

```
} finally {  
    scanner.close();  
}  
table.close();  
admin.close();  
// Drop the table  
admin.disableTable(tablename);  
admin.deleteTable(tablename);  
}  
}
```

# Hbase & Hive Integration

# Any Questions in Your Mind?



# Thank You

- Siva Kumar Bhuchipalli  
Inventor of [www.hadooptutorial.info](http://www.hadooptutorial.info)  
phone no:  
Address:

